



The E/R Model

Prof. Hyuk-Yoon Kwon

<https://sites.google.com/view/seoultech-bigdata>

Most parts are based on slides used in Stanford (<http://web.stanford.edu/class/cs145>)

Contents

■ Review of previous lecture

- Course overview
- Introduction to Databases

■ Today's lecture

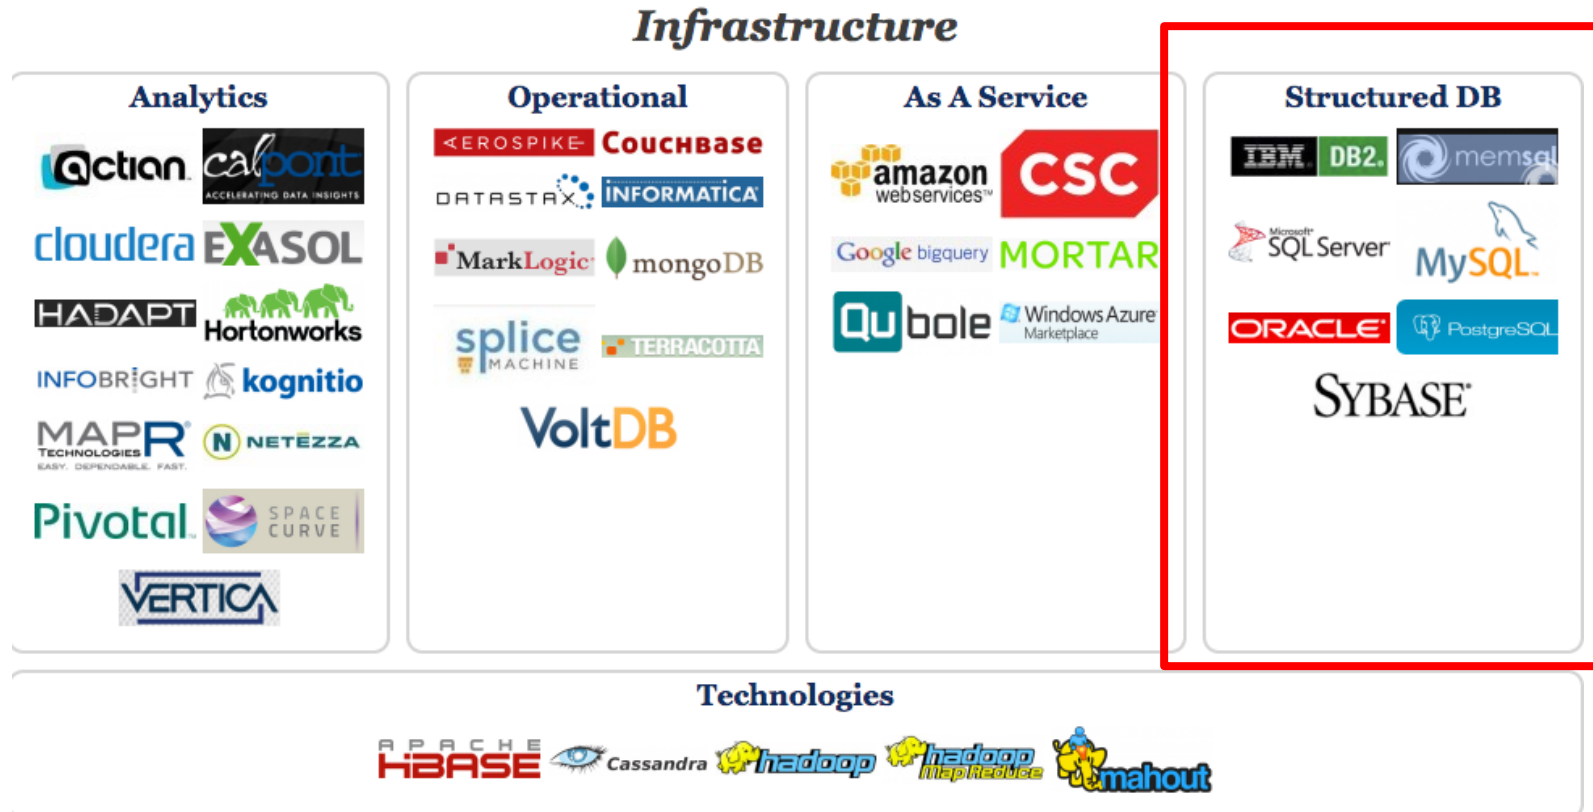
- E/R Basics: Entities & Relations
- E/R Design considerations
- Advanced E/R Concepts



Review of Previous Lecture

Why Data Is Important?

- The world is increasingly driven by data...
- Big Data Landscape...
 - Infrastructure is Changing



The Aim of Course

- 1.** To learn database concepts and basic theory
- 2.** To practice database issues such as SQL and database design
 - Using Oracle Database
- 3.** To make real Web services based on the databases

Course Policy: Grading

■ Exams (70%)

- Mid-term (30%)
- Final (40%)

■ Lab assignments (10%)

■ Individual presentation (20%)

What is a DBMS?

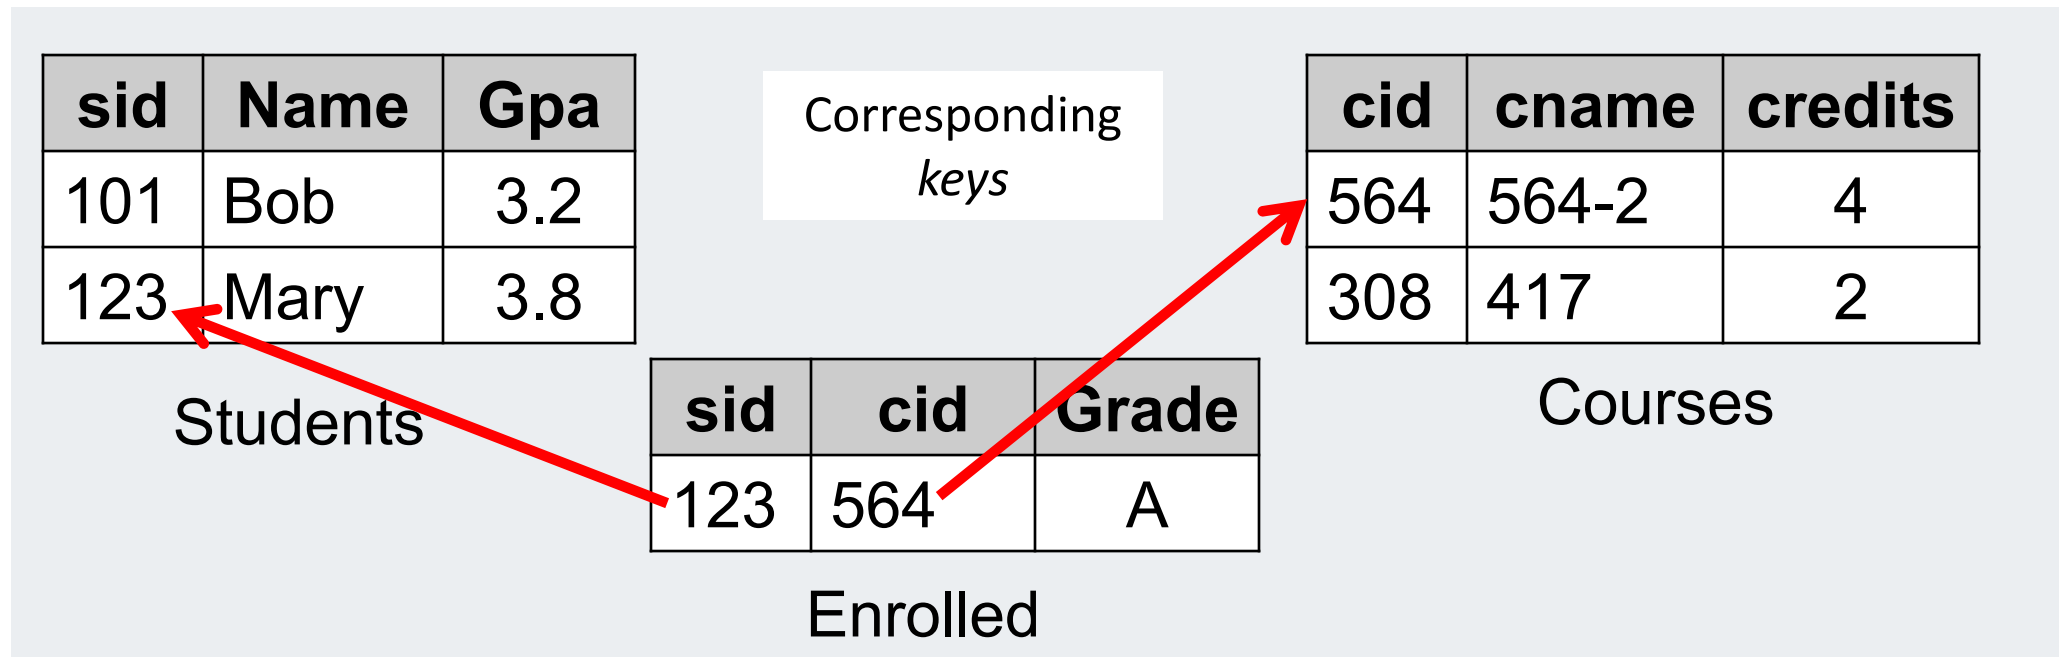
- A large, integrated collection of data
- Models a real-world enterprise
 - *Entities* (e.g., Students, Courses)
 - *Relationships* (e.g., Alice is enrolled in software development and practice)

A **Database Management System (DBMS)** is a piece of software designed to store and manage databases

Modeling the CMS

■ Logical Schema

- Students(*sid: string, name: string, gpa: float*)
- Courses(*cid: string, cname: string, credits: int*)
- Enrolled(*sid: string, cid: string, grade: string*)



Challenges with Many Users

■ Suppose that our CMS application serves 1000's of users or more- what are some challenges?

- Security: Different users, different roles

We won't look at too much in this course, but is extremely important

- Performance: Need to provide concurrent access

Disk/SSD access is slow, DBMS hide the latency by doing more CPU work concurrently

- Consistency: Concurrency can lead to update problems

DBMS allows user to write programs as if they were the **only** user

Transactions

- A key concept is the transaction: an atomic sequence of DB actions (reads/writes)

Atomicity: An action either completes *entirely* or *not at all*

all or nothing



Acct	Balance
a10	20,000
a20	15,000

Transfer \$3k from a10 to a20:

1. Debit \$3k from a10
2. Credit \$3k to a20

*multi-operations
like one action
⇒ atomicity*

Acct	Balance
a10	17,000
a20	18,000

Written naively, in which states is **atomicity** preserved?

- Crash before 1,
- < After 1 but before 2, *⇒ No atomicity*
- After 2.

DB Always preserves atomicity!

Ensuring Atomicity

- DBMS ensures atomicity even if a transaction crashes!
- One way to accomplish this: Write-ahead logging (WAL)
- Key Idea: Keep a log of all the writes done.
 - After a crash, the partially executed transactions are undone using the log

Write-ahead Logging (WAL): Before any action is finalized, a corresponding log entry is forced to disk

All atomicity issues also handled by the DBMS...

Today's Lecture

- 1. E/R Basics: Entities & Relations**
2. E/R Design considerations
3. Advanced E/R Concepts

What you will learn about in this section

1. High-level motivation for the E/R model
2. Entities
3. Relations

Database Design

■ Database design: Why do we need it?

- Agree on structure of the database before deciding on a particular implementation

DB 구조 한번 만든번 수정하기 힘들다. 구조 잘 만들기 하면 더욱적 사용 가능. 효율적!

■ Consider issues such as:

- What entities to model
- How entities are related
- What constraints exist in the domain *→ Student ID 는 숫자다, 등 ...*
- How to achieve good designs *→ No correct answer
find better design!
removing redundancy ⇒ efficient model*

■ Several formalisms exist

- We discuss one flavor of E/R diagrams

Entities/Relationship diagram

Database Design Process

1. Requirements Analysis

2. Conceptual Design

3. Logical, Physical, Security, etc.

1. Requirements analysis

often from non-technical people → little knowledge about DB
Requirements of target application or service

Technical and non-technical people are involved

- What is going to be stored?
- How is it going to be used? (*update period, making index, ...*)
- What are we going to do with the data? *- purpose*
- Who should access the data? *≠ roles of users*
es. → separate private data from public, open data
↳ only for users with authorization



Database Design Process

1. Requirements Analysis

2. Conceptual Design

3. Logical, Physical, Security, etc.

2. Conceptual Design

- A high-level description of the database
- Sufficiently precise that technical people can understand it
- But, not so precise that non-technical people can't participate

↳ abstract

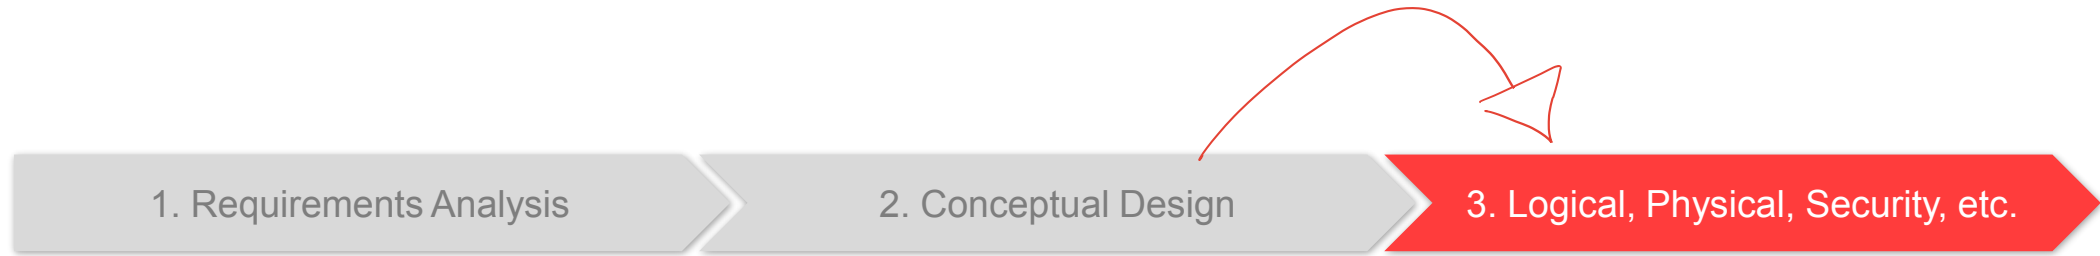
E/R \Rightarrow actual model for build DB

Common tool is needed to communicate

This is where **E/R** fits in.

E/R \Rightarrow bridge between tech / non-tech

Database Design Process



3. More:

- Logical Database Design *⇒ identify Entities & relationship*
- Physical Database Design *⇒ how often update, retrieve / index / where to store ⇒ [improving efficiency of DB]*
- Security Design *⇒ different user roles*
 - └ differentiate roles in multiple level*
 - └ assign different authorities*

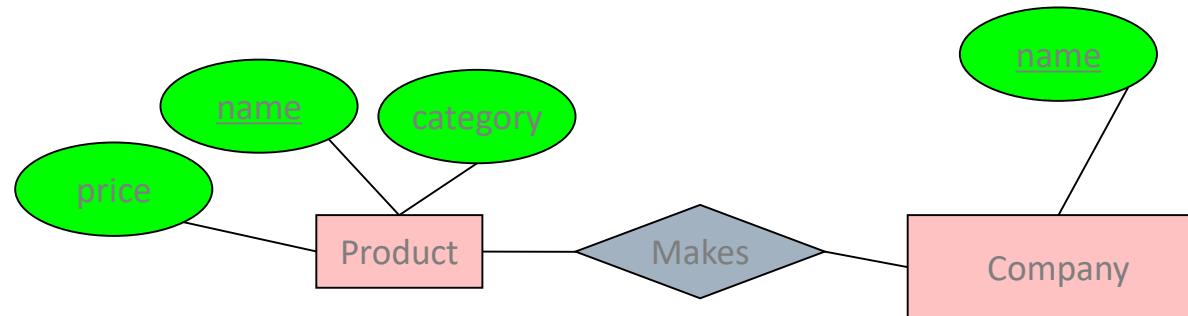
Database Design Process

1. Requirements Analysis

2. Conceptual Design

3. Logical, Physical, Security, etc.

E/R Model & Diagrams used



This process
is iterated ^{반복하다}
many times

until final model

E/R is a *visual syntax* for DB design which is ***precise enough*** for technical points, but ***abstracted enough*** for non-technical people

Impact of the ER model

- **The E/R model is one of the most cited articles in Computer Science**

- *“The Entity-Relationship model – toward a unified view of data”* Peter Chen, 1976

- **Used by companies big and small**

- You'll know it soon enough

Entities and Entity Sets

■ Entities & entity sets are the primitive unit of the E/R model

- Entities are the individual objects, which are members of entity sets

- Ex: A specific person or product

actual data

- Entity sets are the *classes* or *types* of objects in our model

- Entity sets represent the sets of all possible entities

- Ex: Person, Product

- These are what is shown in E/R diagrams - as rectangles

actual data type

Product

Person

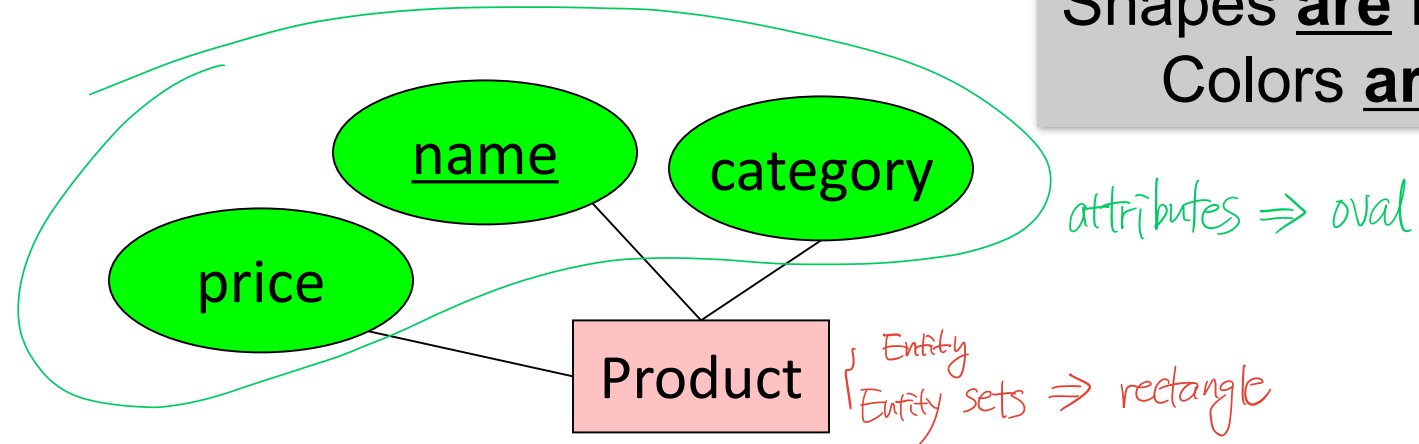
These represent entity sets

rectangle

Entities and Entity Sets

■ An entity set has attributes

- Represented by ovals attached to an entity set

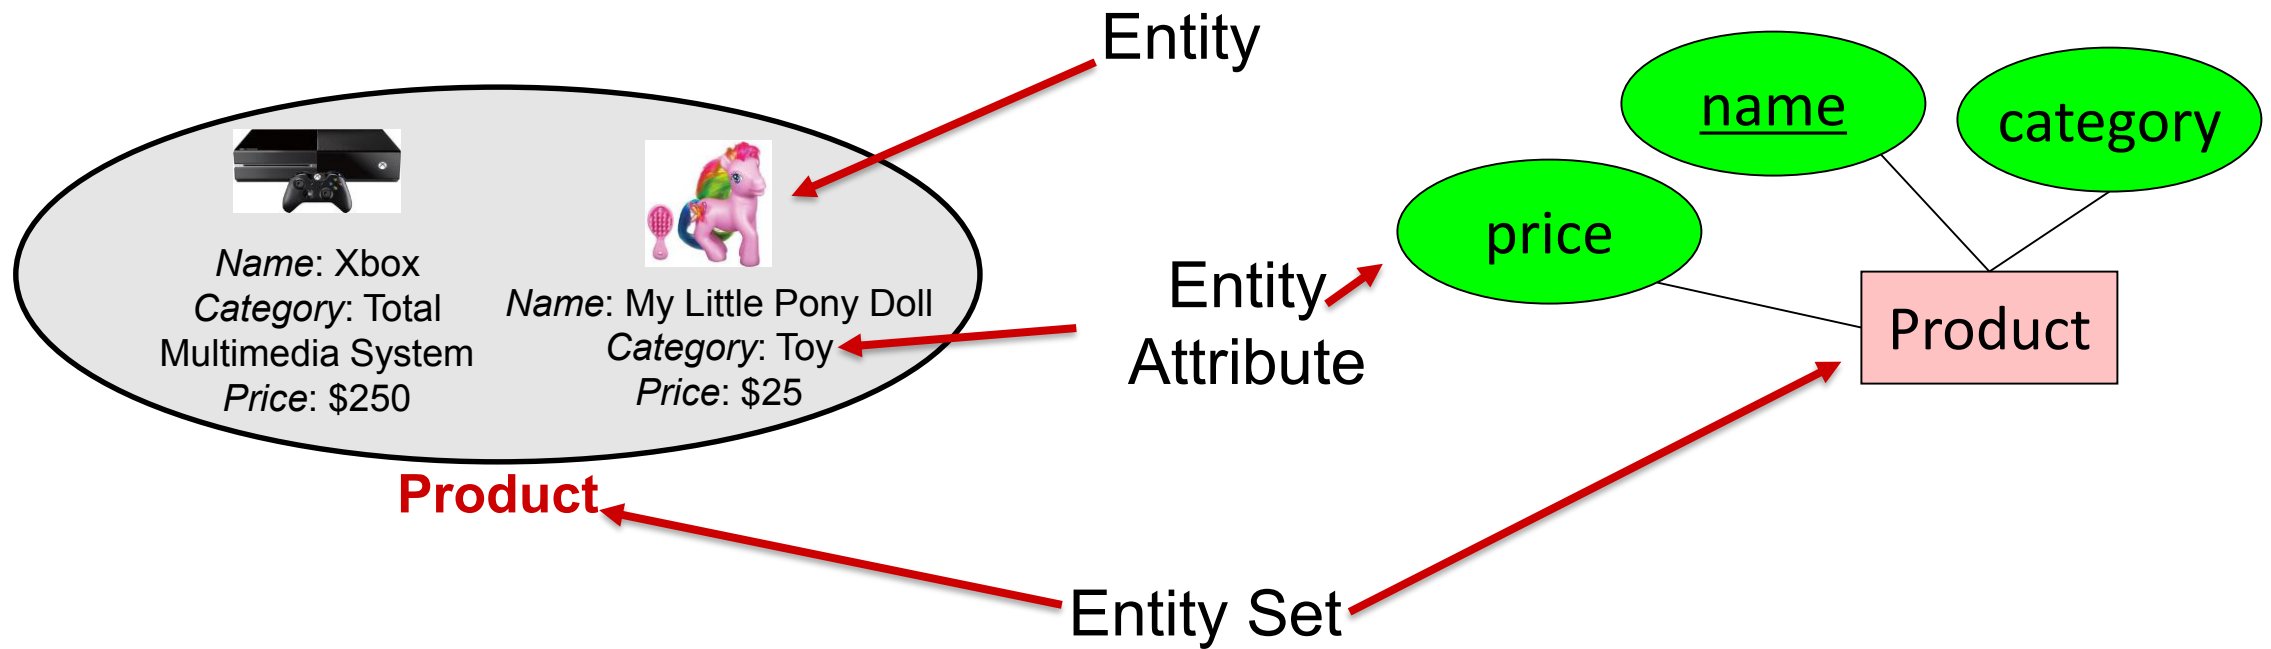


Shapes are important.
Colors are not.

Entities vs. Entity Sets

Example:

Entities are not explicitly represented in E/R diagrams!



Keys


- A key is a **minimal set** of attributes **that uniquely identifies** an entity.

각 Entity의 최소구분능력.

{학번 + 이름} \Rightarrow key 아님. 학번만으로 구분가능.

{상품명 + 카테고리} \Rightarrow 상품명 중복이 있다면, key이다. 등 외는 있어야 구분하기 때문.

Denote elements of the primary key by underlining.



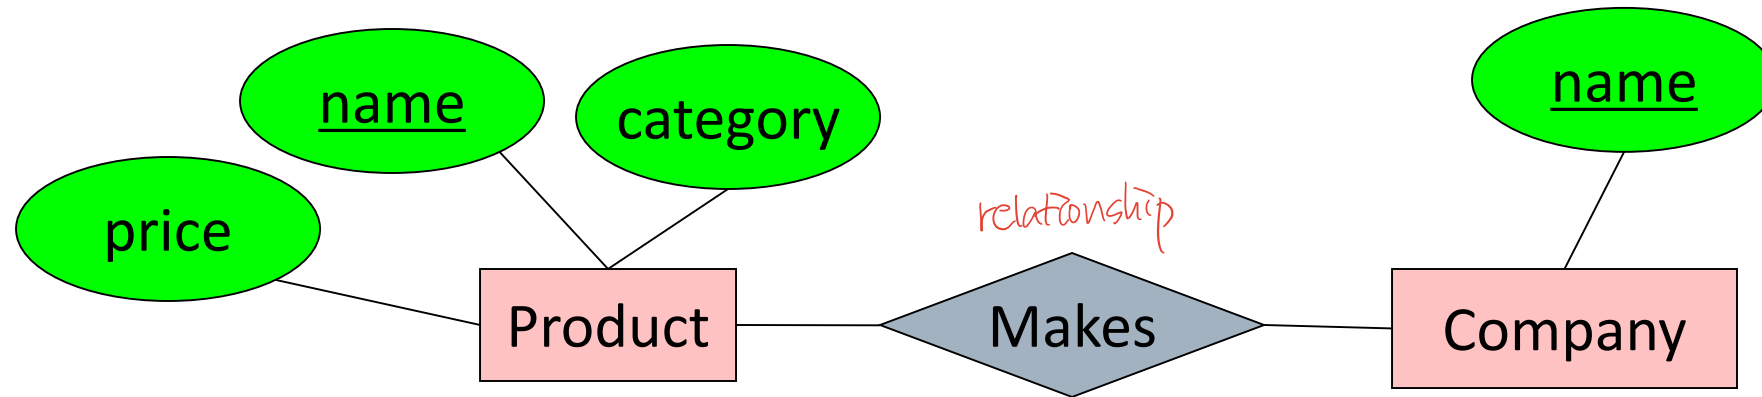
Here, {name, category} is **not** a key (it is not *minimal*).

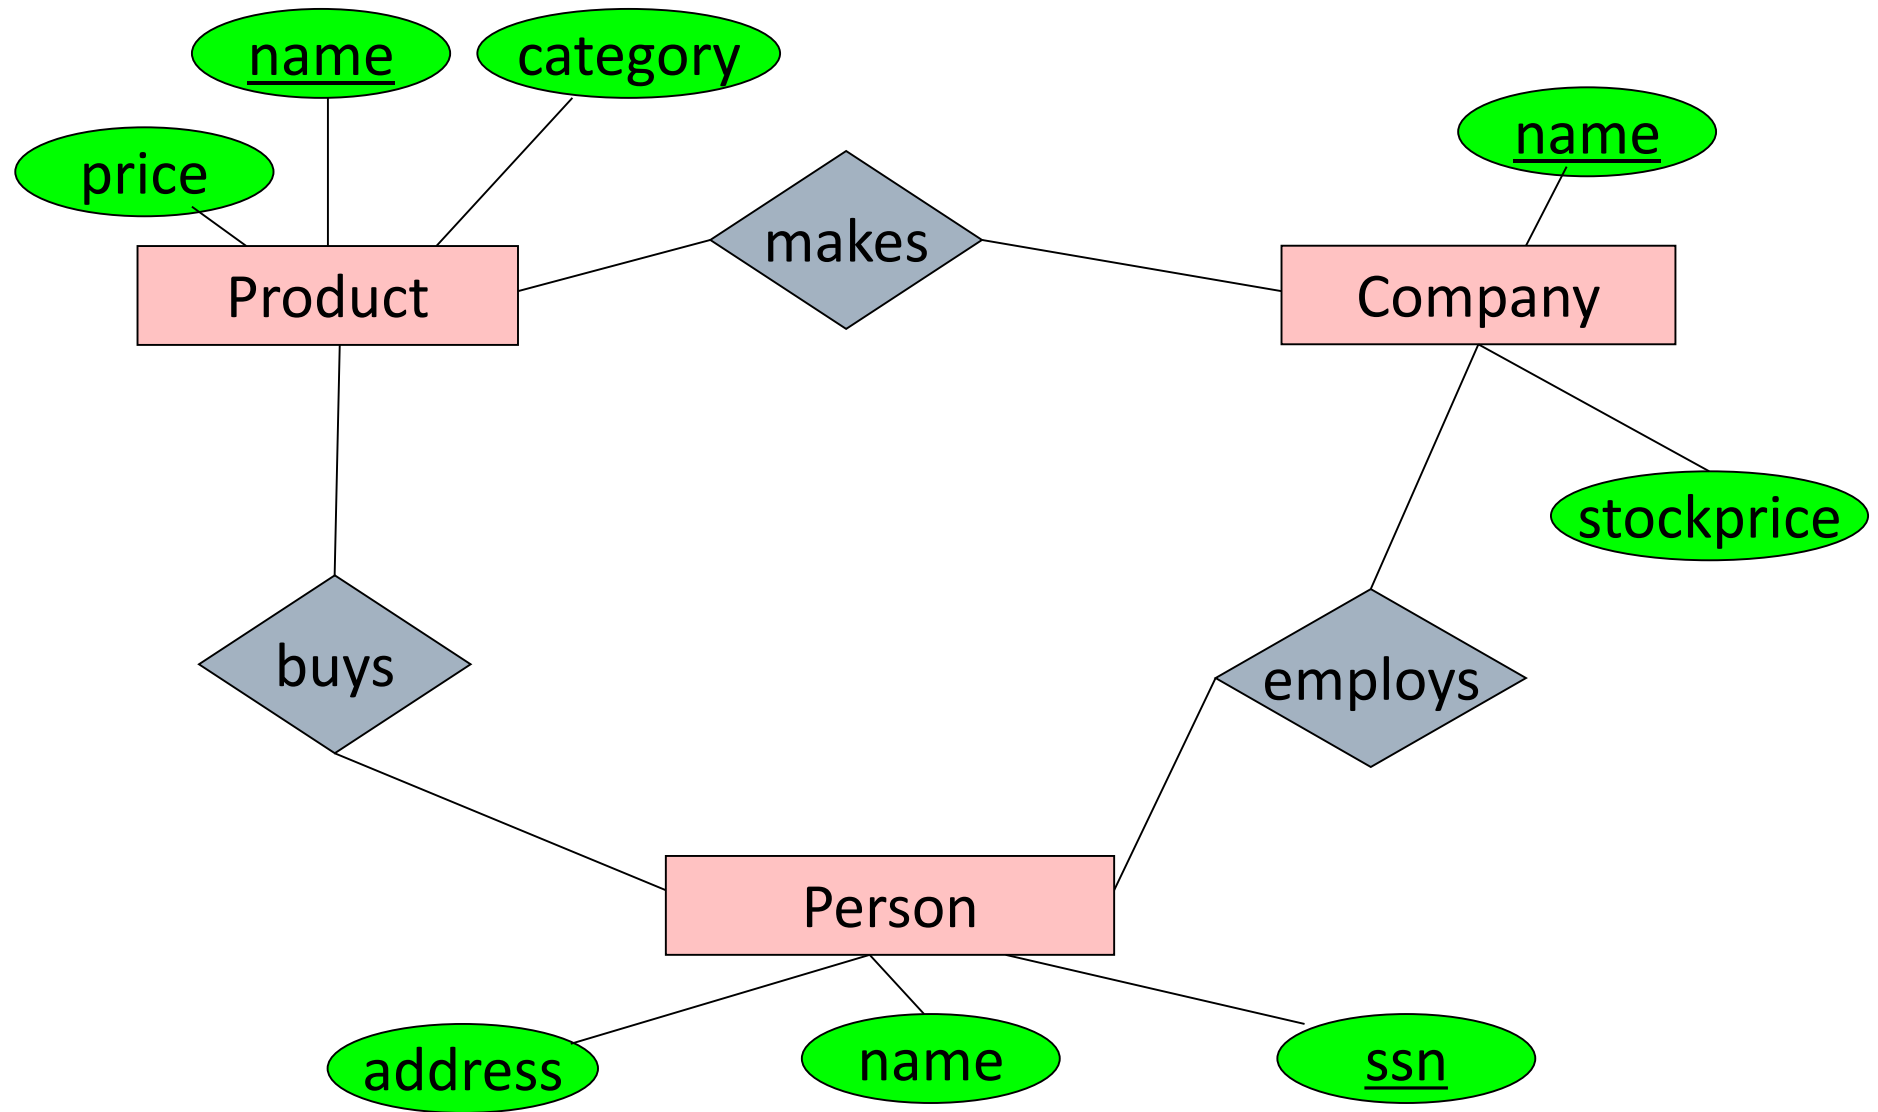
If it were, what would it mean?

The E/R model forces us to designate a single **primary** key, though there may be multiple candidate keys

The R in E/R: Relationships

- A relationship is between two entities






What is a Relationship?

■ *A mathematical definition:*


- Let A, B be sets
 - $A=\{1,2,3\}$, $B=\{a,b,c,d\}$



What is a Relationship?

■ A mathematical definition:


- Let A, B be sets
 - $A=\{1,2,3\}$, $B=\{a,b,c,d\}$
- $A \times B$ (the **cross-product**) is the set of all pairs (a,b)
 - $A \times B = \{(1,a), (1,b), (1,c), (1,d), (2,a), (2,b), (2,c), (2,d), (3,a), (3,b), (3,c), (3,d)\}$



What is a Relationship?

■ A mathematical definition:

- Let A, B be sets
 - $A=\{1,2,3\}$, $B=\{a,b,c,d\}$,
- $A \times B$ (the **cross-product**) is the set of all pairs (a,b)
 - $A \times B = \{(1,a), (1,b), (1,c), (1,d), (2,a), (2,b), (2,c), (2,d), (3,a), (3,b), (3,c), (3,d)\}$
- We define a relationship to be a subset of $A \times B$
 - $R = \{(1,a), (2,c), (2,d), (3,b)\}$





What is a Relationship?


■ A mathematical definition:

- Let A, B be sets
- $A \times B$ (the **cross-product**) is the set of all pairs
- A relationship is a subset of $A \times B$

■ Makes is relationship- it is a *subset* of $\text{Product} \times \text{Company}$:



What is a Relationship?



A relationship between entity sets **P** and **C** is a ***subset of all possible pairs of entities in P and C,*** with tuples uniquely identified by ***P and C's keys***


What is a Relationship?

Company

<u>name</u>
GizmoWorks
GadgetCorp

Product

<u>name</u>	category	price
Gizmo	Electronics	\$9.99
GizmoLite	Electronics	\$7.50
Gadget	Toys	\$5.50



A relationship between entity sets P and C is a *subset of all possible pairs of entities in P and C* , with tuples uniquely identified by P and C 's keys


What is a Relationship?

Company

<u>name</u>
GizmoWorks
GadgetCorp


Product

<u>name</u>	category	price
Gizmo	Electronics	\$9.99
GizmoLite	Electronics	\$7.50
Gadget	Toys	\$5.50



Company C × Product P

<u>C.name</u>	<u>P.name</u>	P.category	P.price
GizmoWorks	Gizmo	Electronics	\$9.99
GizmoWorks	GizmoLite	Electronics	\$7.50
GizmoWorks	Gadget	Toys	\$5.50
GadgetCorp	Gizmo	Electronics	\$9.99
GadgetCorp	GizmoLite	Electronics	\$7.50
GadgetCorp	Gadget	Toys	\$5.50



A relationship between entity sets P and C is a *subset of all possible pairs of entities in P and C*, with tuples uniquely identified by *P and C's keys*


What is a Relationship?

Company

<u>name</u>
GizmoWorks
GadgetCorp

Product

<u>name</u>	category	price
Gizmo	Electronics	\$9.99
GizmoLite	Electronics	\$7.50
Gadget	Toys	\$5.50




Company C × Product P

<u>C.name</u>	<u>P.name</u>	P.category	P.price
GizmoWorks	Gizmo	Electronics	\$9.99
GizmoWorks	GizmoLite	Electronics	\$7.50
GizmoWorks	Gadget	Toys	\$5.50
GadgetCorp	Gizmo	Electronics	\$9.99
GadgetCorp	GizmoLite	Electronics	\$7.50
GadgetCorp	Gadget	Toys	\$5.50



Makes

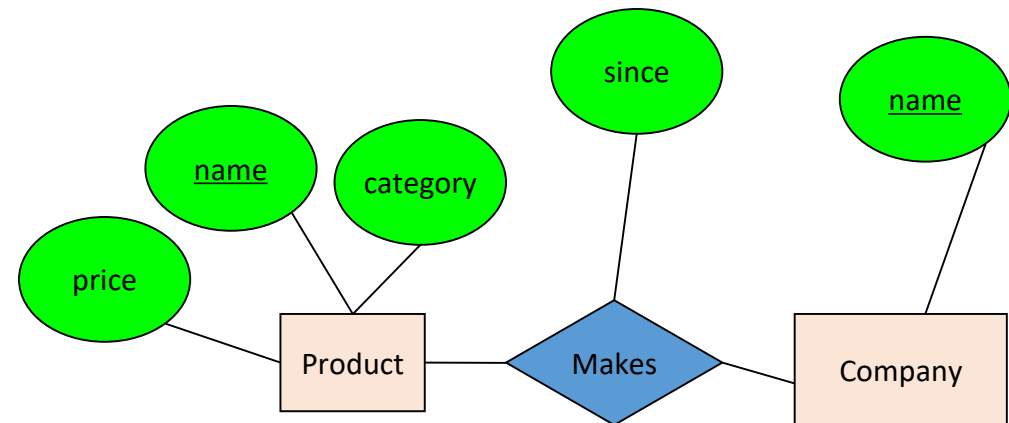
<u>C.name</u>	<u>P.name</u>
GizmoWorks	Gizmo
GizmoWorks	GizmoLite
GadgetCorp	Gadget



A relationship between entity sets P and C is a *subset of all possible pairs of entities in P and C*, with tuples uniquely identified by *P and C's keys*

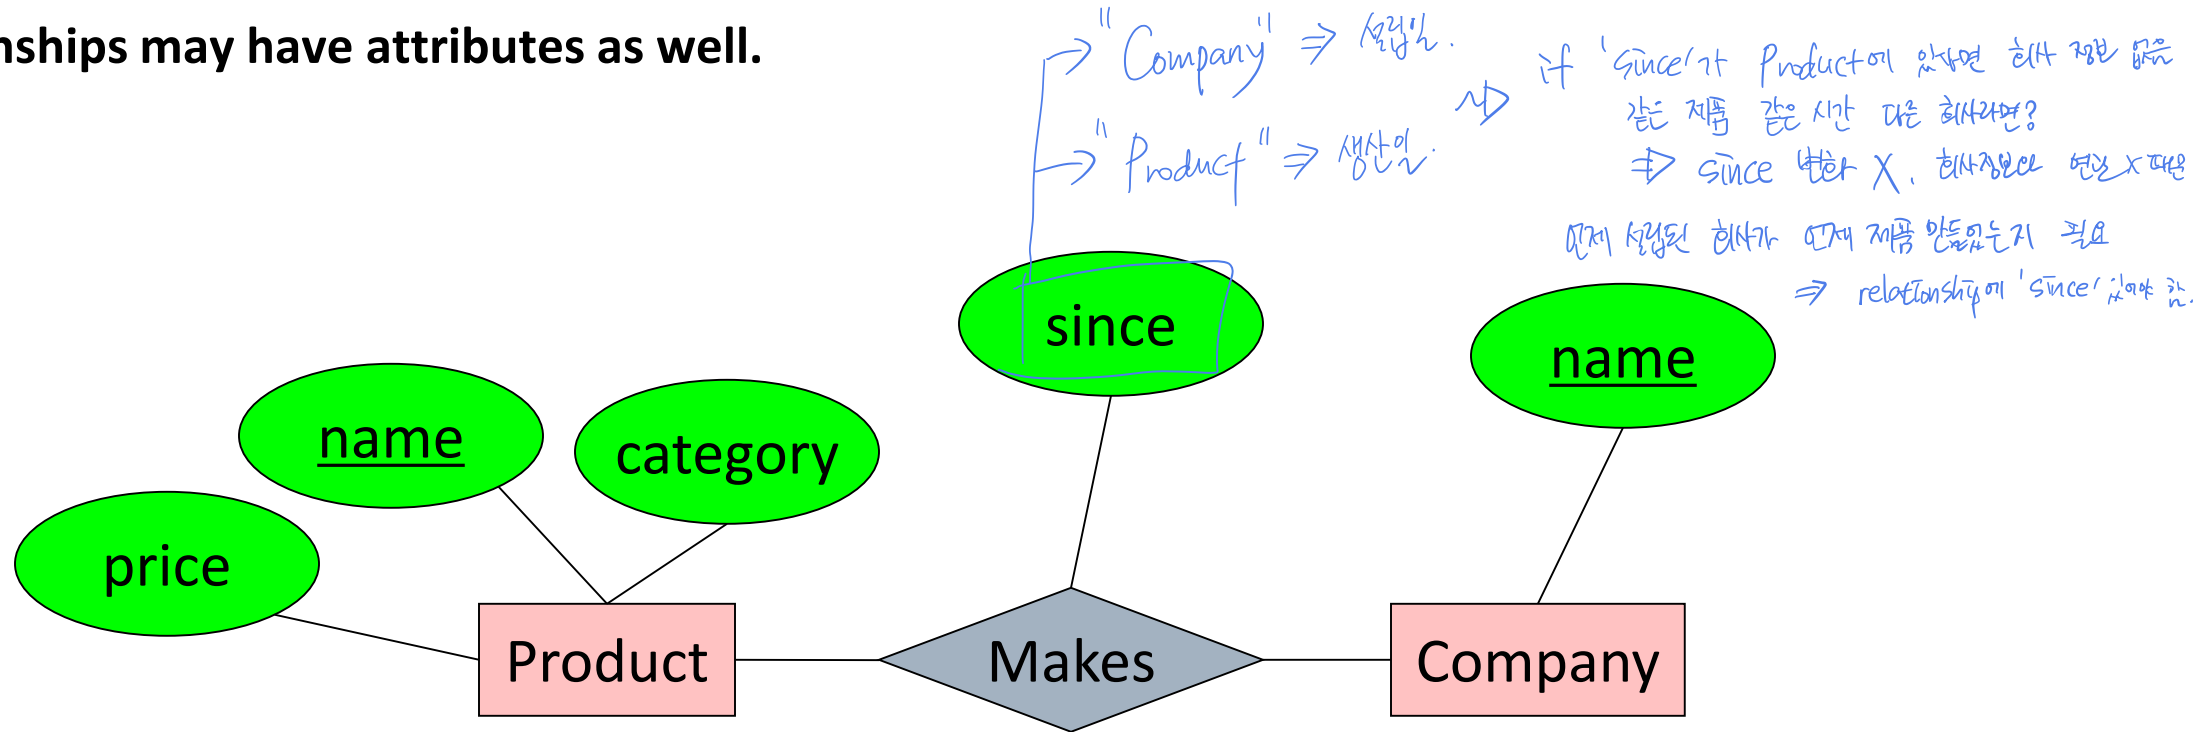
What is a Relationship?

- There can only be one relationship for every unique combination of entities
- This also means that the relationship is uniquely determined by the keys of its entities
- *Example: the “key” for Makes (to right) is {Product.name, Company.name}*



Relationships and Attributes

- Relationships may have attributes as well.



For example: "since" records when company started making a product

Note: "since" is implicitly unique per pair here! Why?

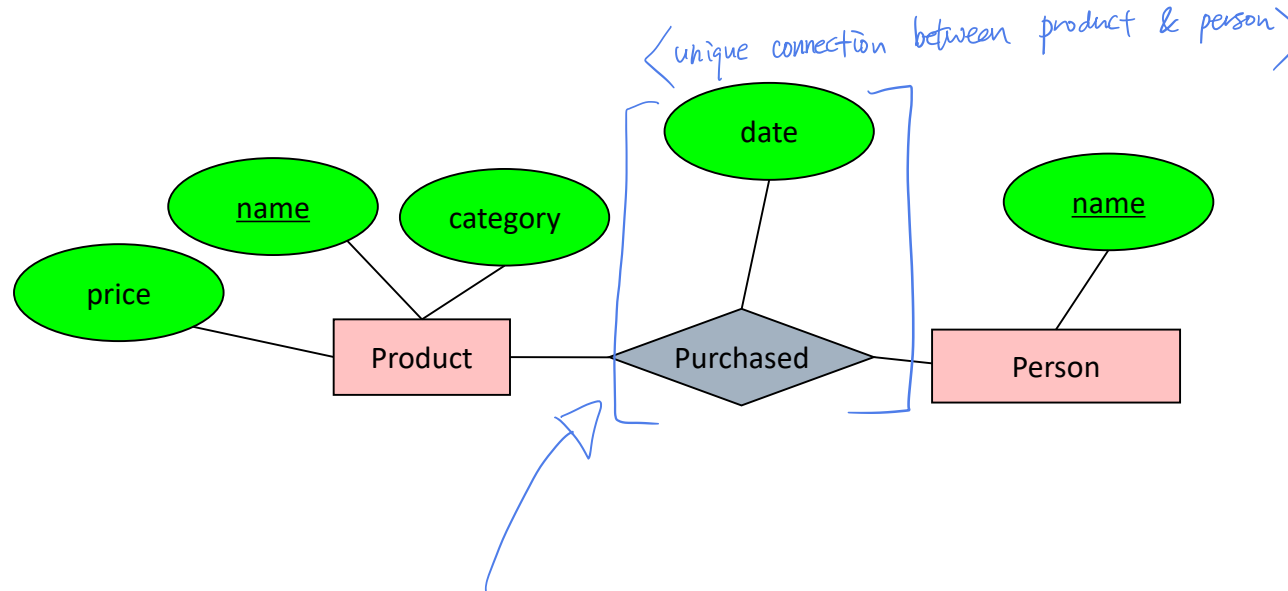
represented by relationship between 'Product' & 'Company'

Note #2: Why not "how long"?

since ⇒ able to calculate ⇒ 0
how long ⇒ update requires ⇒ X
⇒ efficient! maintenance cost!

Decision: Relationship vs. Entity?

■ Q: What does this say?

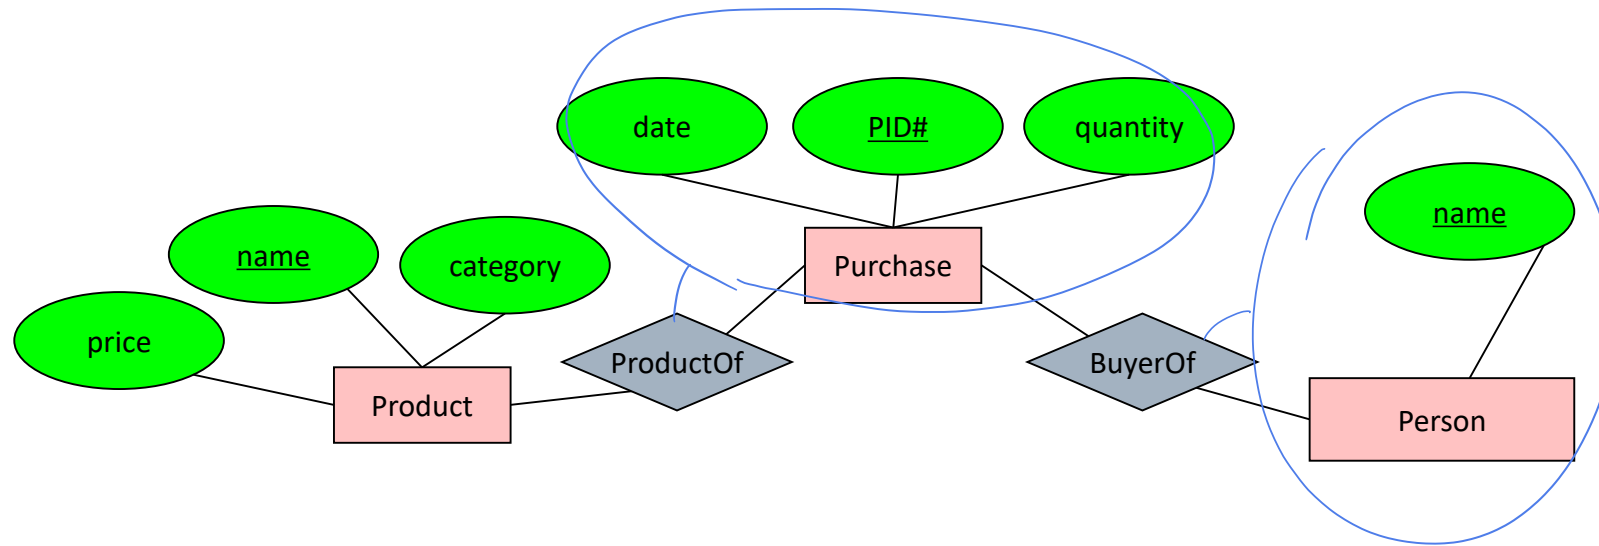


■ A: A person can only buy a specific product once (on one date)

Modeling something as a relationship makes it unique;
what if not appropriate?

Decision: Relationship vs. Entity?

■ What about this way?



■ *Now we can have multiple purchases per product, person pair!*

We can always use a **new entity** instead of a relationship. For example, to permit multiple instances of each entity combination!



Practice1

Draw an E/R diagram for football

Use the following simplified model of a football season
(concepts to include are underlined):




Teams play each other in Games.

Each pair of teams can play each other multiple times



Players belong to Teams (assume no trades / changes).



A Game is made up of Plays that result in a yardage gain/loss, and potentially a touchdown



A Play will contain either a Pass from one player to another, or a Run by one player

Today's Lecture



1. E/R Basics: Entities & Relations
- 2. E/R Design considerations**
3. Advanced E/R Concepts

What you will learn about in this section

- 1. Relationships cont'd: multiplicity, multi-way**
- 2. Design considerations**
- 3. Conversion to relational schema**



Multiplicity of E/R Relationships

One-to-one:





Indicated using
arrows



Many-to-one:



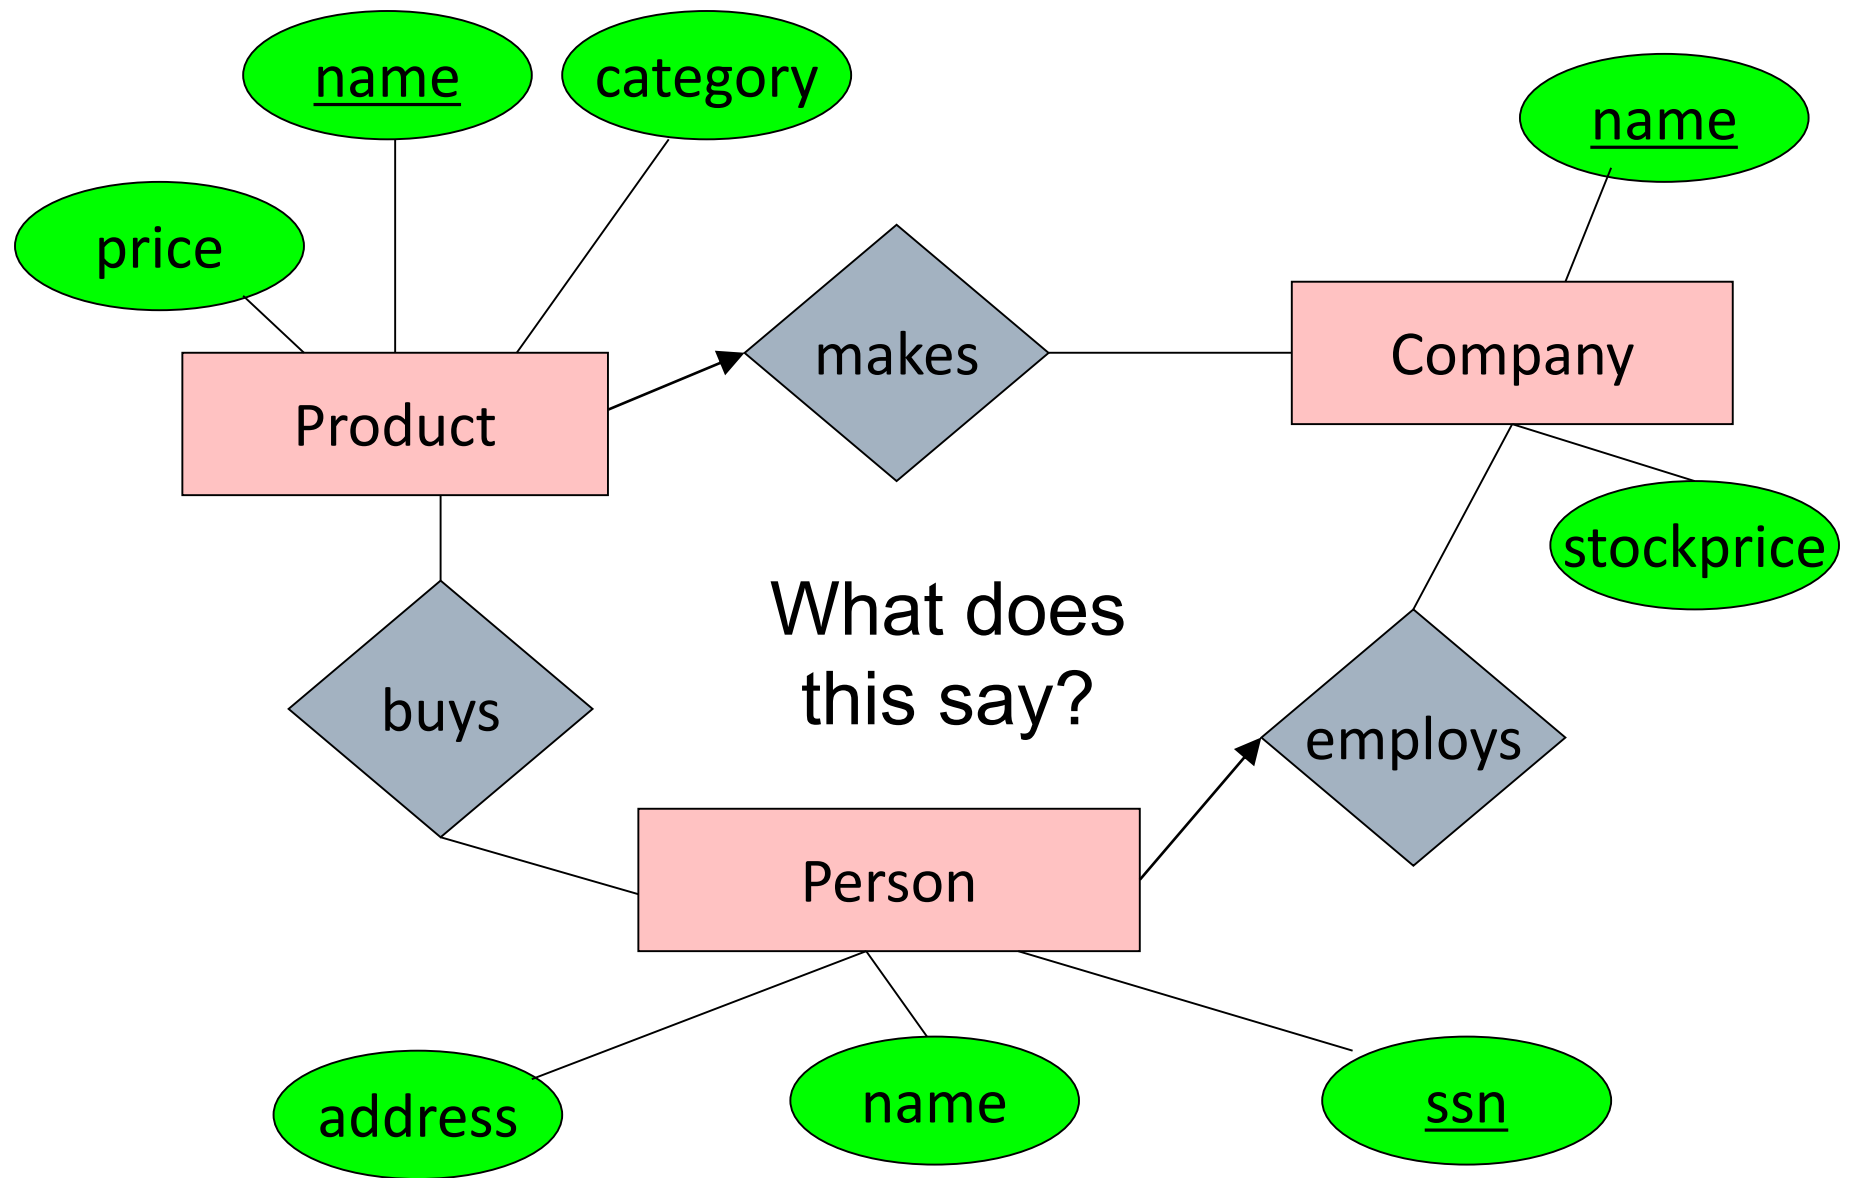
One-to-many:



Many-to-many:




$X \rightarrow Y$ means
there exists a
function mapping
from X to Y (*recall*
the definition of a
function)




Multi-way Relationships

- How do we model a purchase relationship between buyers, products and stores?




Arrows in Multiway Relationships

Q: What does the arrow mean ?

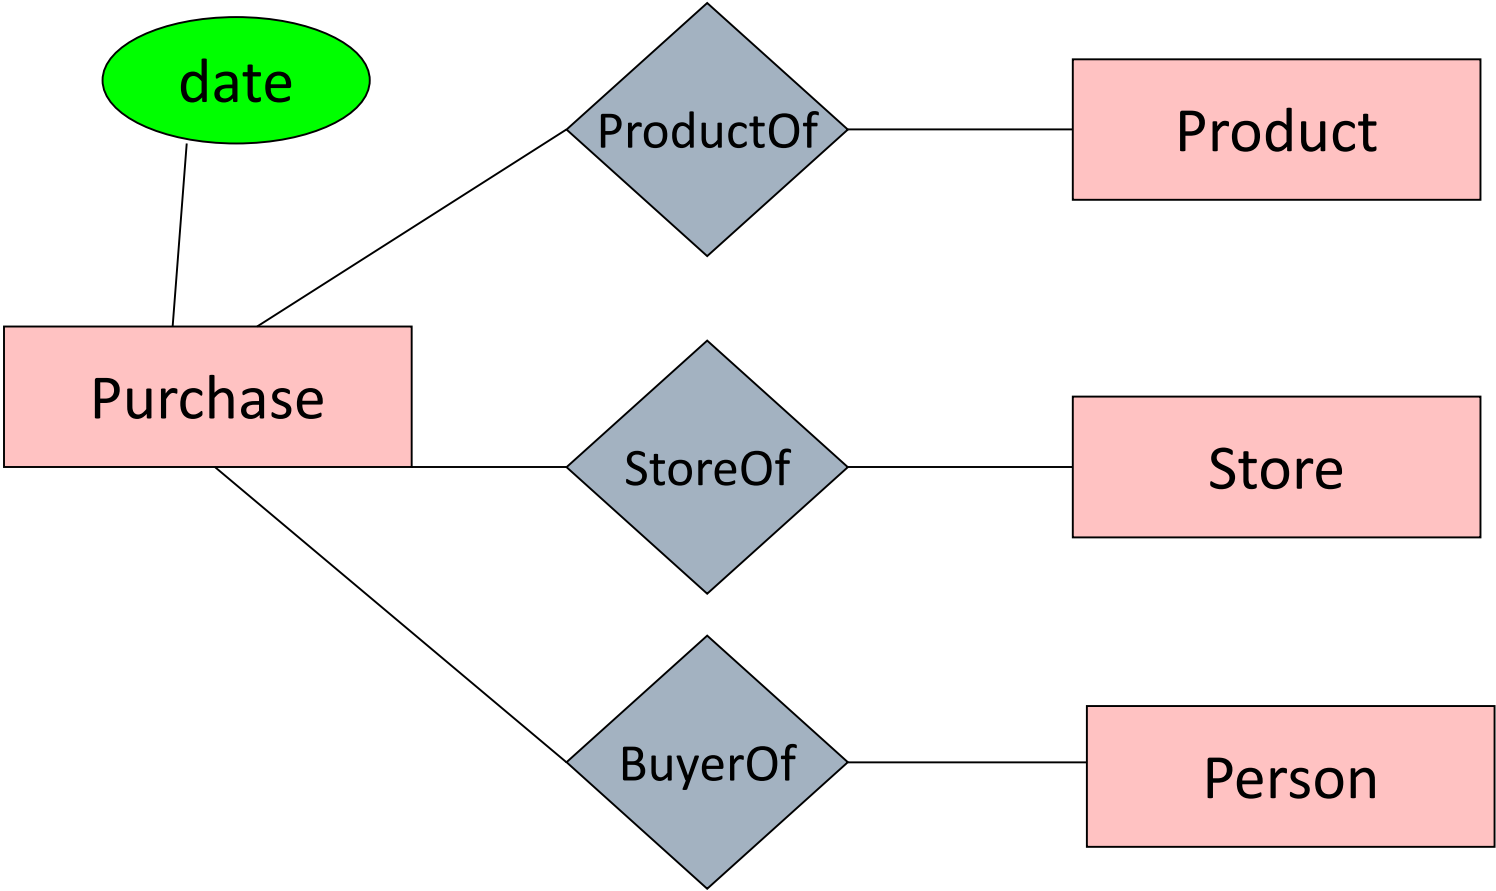


Arrows in Multiway Relationships

Q: What does the arrow mean ?




Converting Multi-way Relationships to Binary




From what we had on previous slide to this - what did we do?

Decision: Multi-way or New Entity + Binary?

Multi-way Relationship




Entity + Binary




Should we use a single **multi-way relationship** or a ***new entity with binary relations?***

Decision: Multi-way or New Entity + Binary?

(A) Multi-way Relationship



(B) Entity + Binary




Multiple purchases per
(product, store, person)
combo possible here!

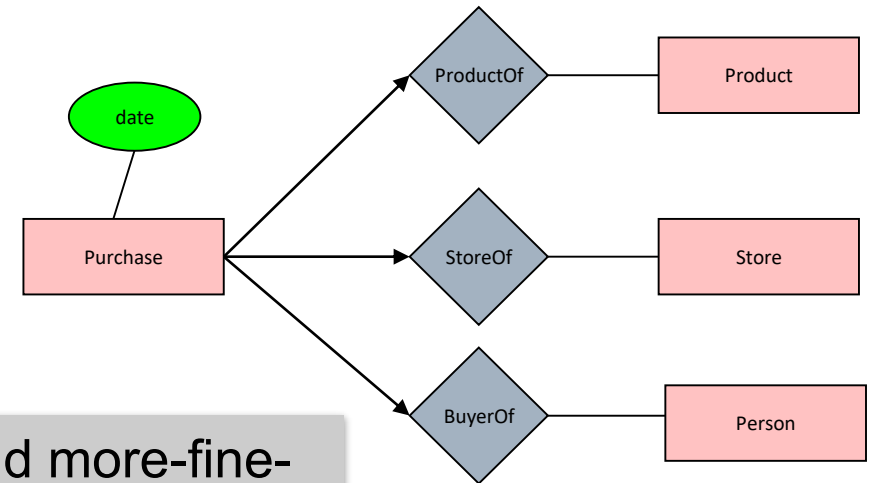
- *Covered earlier:* (B) is useful if we want to have multiple instances of the “relationship” per entity combination

Decision: Multi-way or New Entity + Binary?

(A) Multi-way Relationship



(B) Entity + Binary




We can add more-fine-grained constraints here!


- (B) is also useful when we want to add details (constraints or attributes) to the relationship
 - “A person who shops in only one store”
 - “How long a person has been shopping at a store”

Decision: Multi-way or New Entity + Binary?

(A) Multi-way Relationship





(B) Entity + Binary



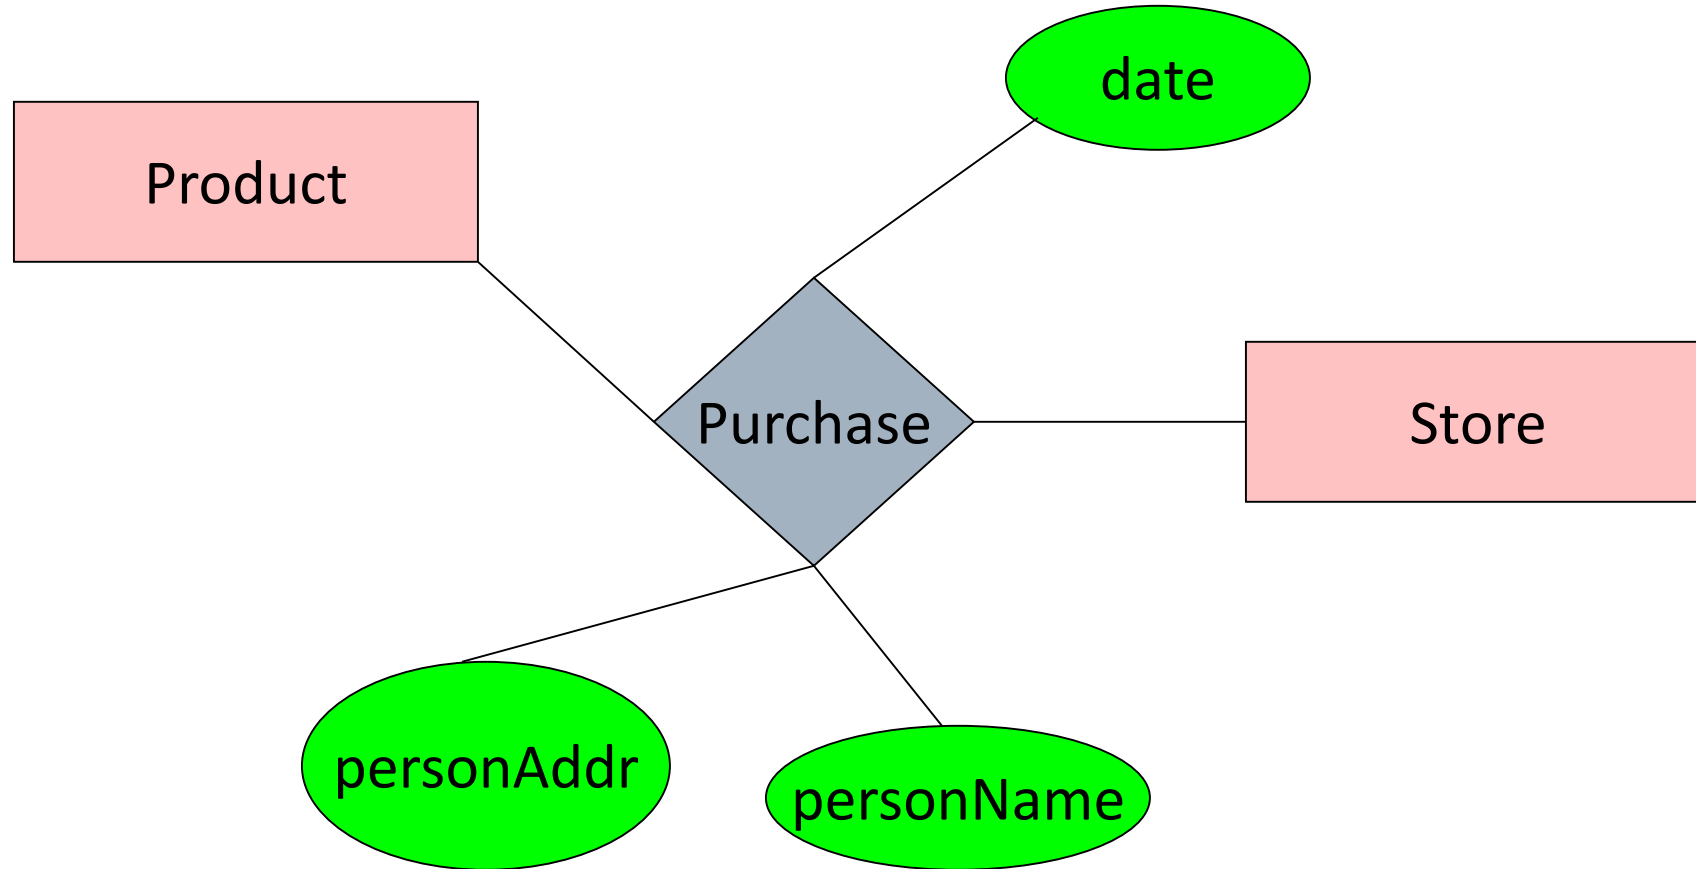
- (A) is useful when a relationship really is between multiple entities
 - *Ex: A three-party legal contract*

Design Principles

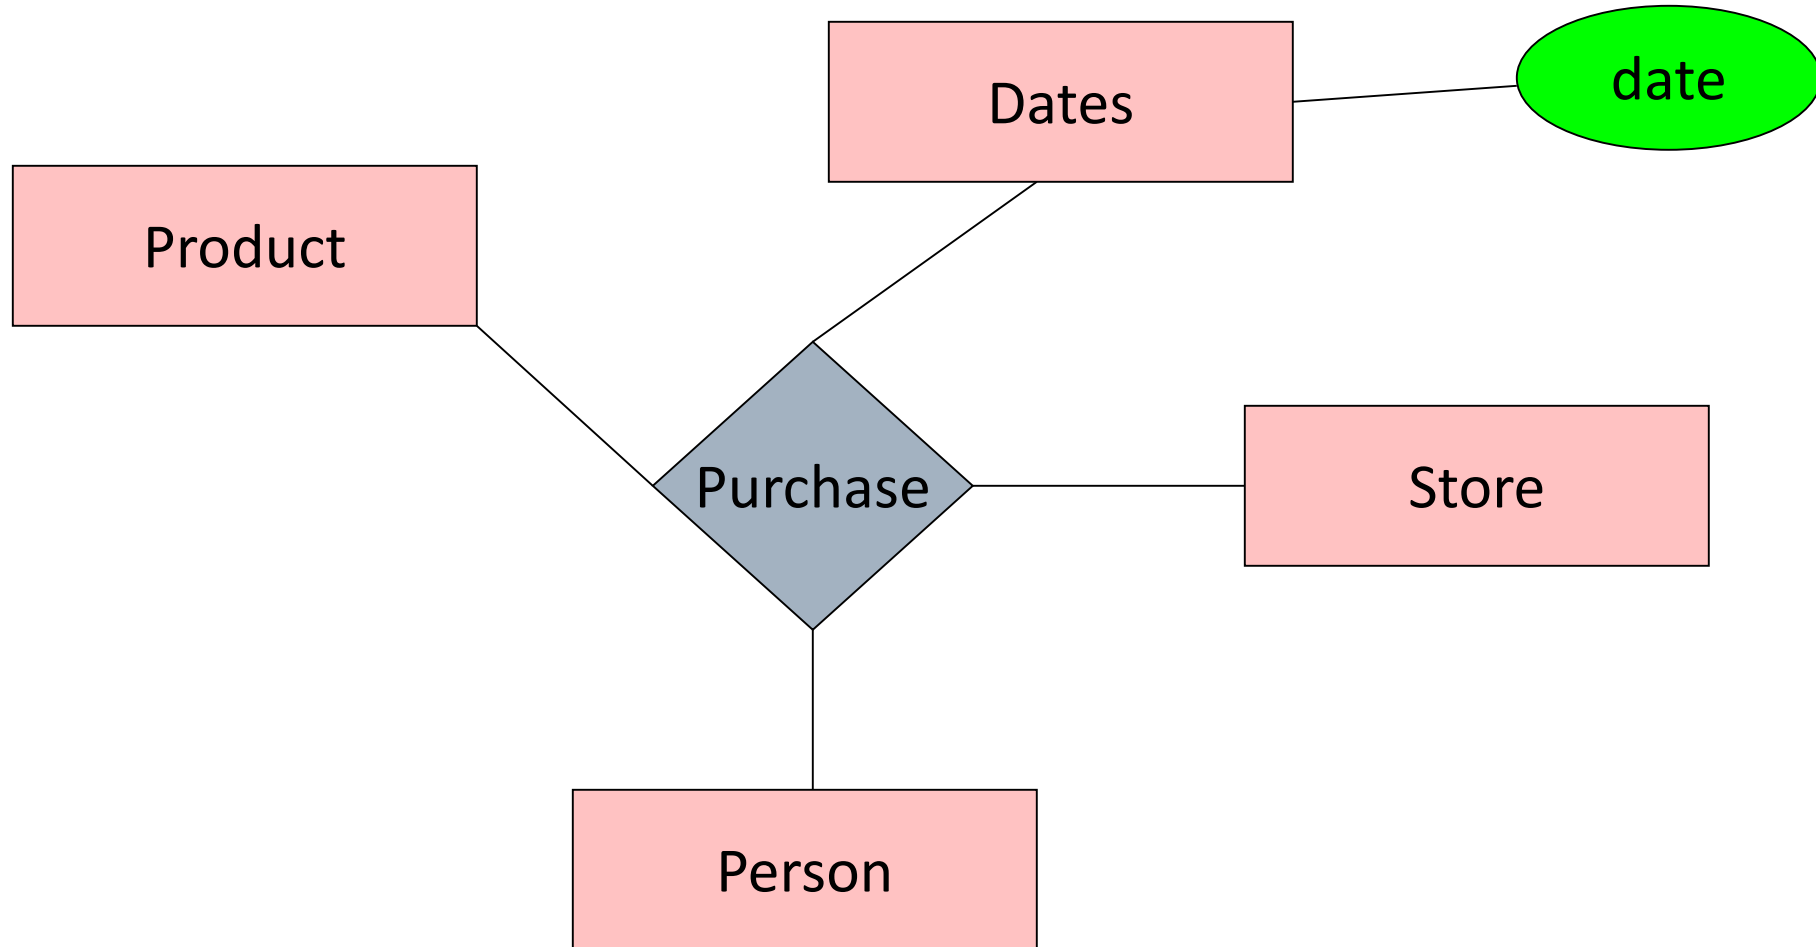
What's wrong with these examples?



Design Principles: What's Wrong?




Design Principles: What's Wrong?




Examples: Entity vs. Attribute

Should address
(A) be an
attribute?




Or (B) be an
entity?



Examples: Entity vs. Attribute

Should address
(A) be an
attribute?




How do we handle
employees with multiple
addresses here?


How do we handle
addresses where internal
structure of the address
(e.g. zip code, state) is
useful?

Examples: Entity vs. Attribute

Should address
(A) be an
attribute?



Or (B) be an
entity?



In general, when we want to record several
values, we choose new entity


From E/R Diagrams to Relational Schema

Key concept:

Both *Entity sets* and *Relationships* become relations (tables in RDBMS)

From E/R Diagrams to Relational Schema

- An entity set becomes a relation (multiset of tuples / table)
 - Each tuple is one entity
 - Each tuple is composed of the entity's attributes, and has the same primary key




Product

<u>name</u>	price	category
Gizmo1	99.99	Camera
Gizmo2	19.99	Edible

From E/R Diagrams to Relational Schema

```
CREATE TABLE Product(  
  name    CHAR(50) PRIMARY KEY,  
  price   DOUBLE,  
  category VARCHAR(30)  
)
```

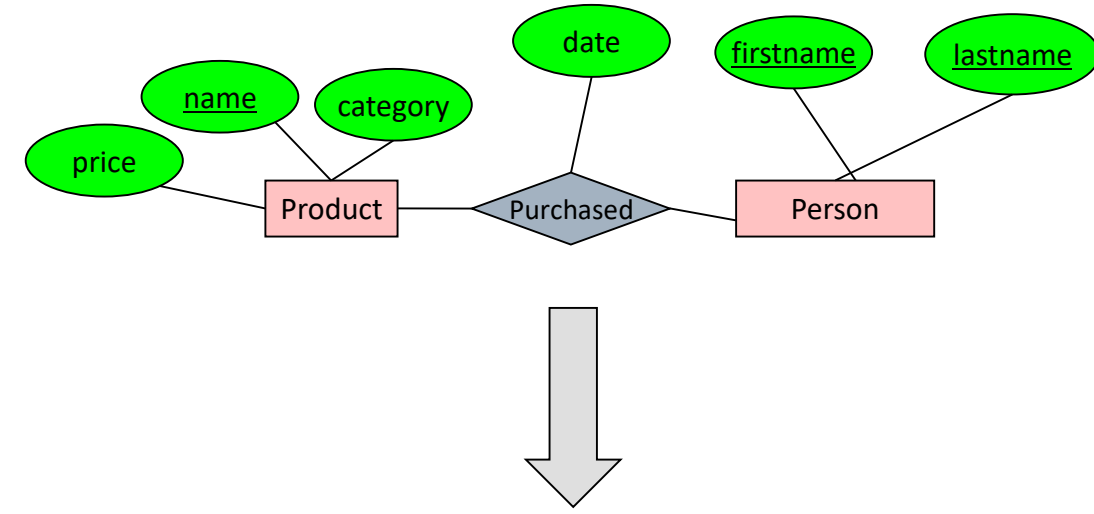


Product

<u>name</u>	price	category
Gizmo1	99.99	Camera
Gizmo2	19.99	Edible

From E/R Diagrams to Relational Schema

- A relation between entity sets A_1, \dots, A_N *also* becomes a multiset of tuples / a table
 - Each row/tuple is one relation, i.e. one unique combination of entities (a_1, \dots, a_N)
 - Each row/tuple is
 - composed of the **union of the entity sets' keys**
 - has the entities' primary keys as foreign keys
 - has the union of the entity sets' keys as primary key

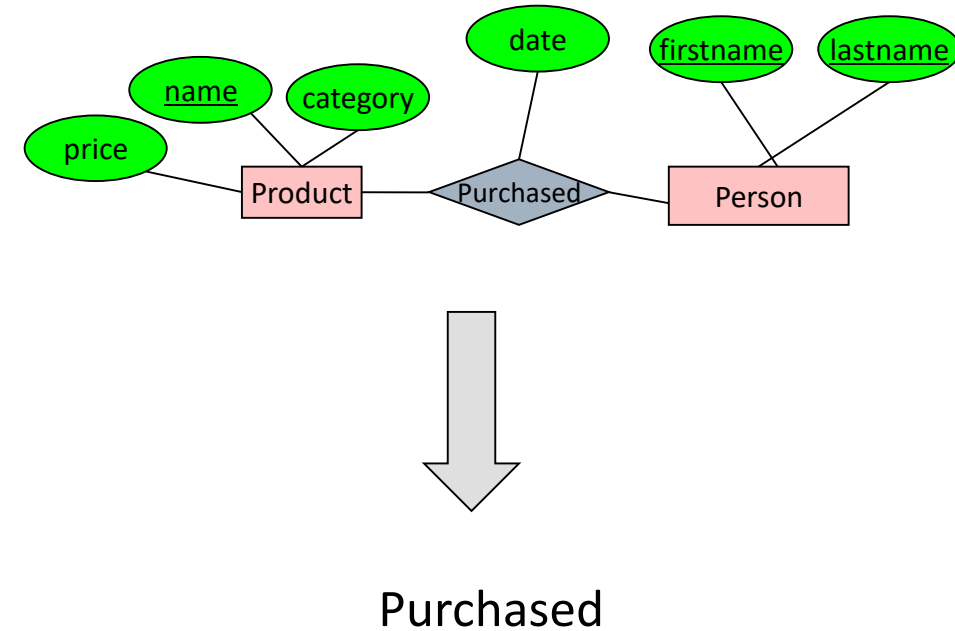


Purchased

<u>name</u>	<u>firstname</u>	<u>lastname</u>	<u>date</u>
Gizmo1	Bob	Alice	01/01/15
Gizmo2	Alice	Bob	01/03/15
Gizmo1	Joe	Smith	01/05/15

From E/R Diagrams to Relational Schema

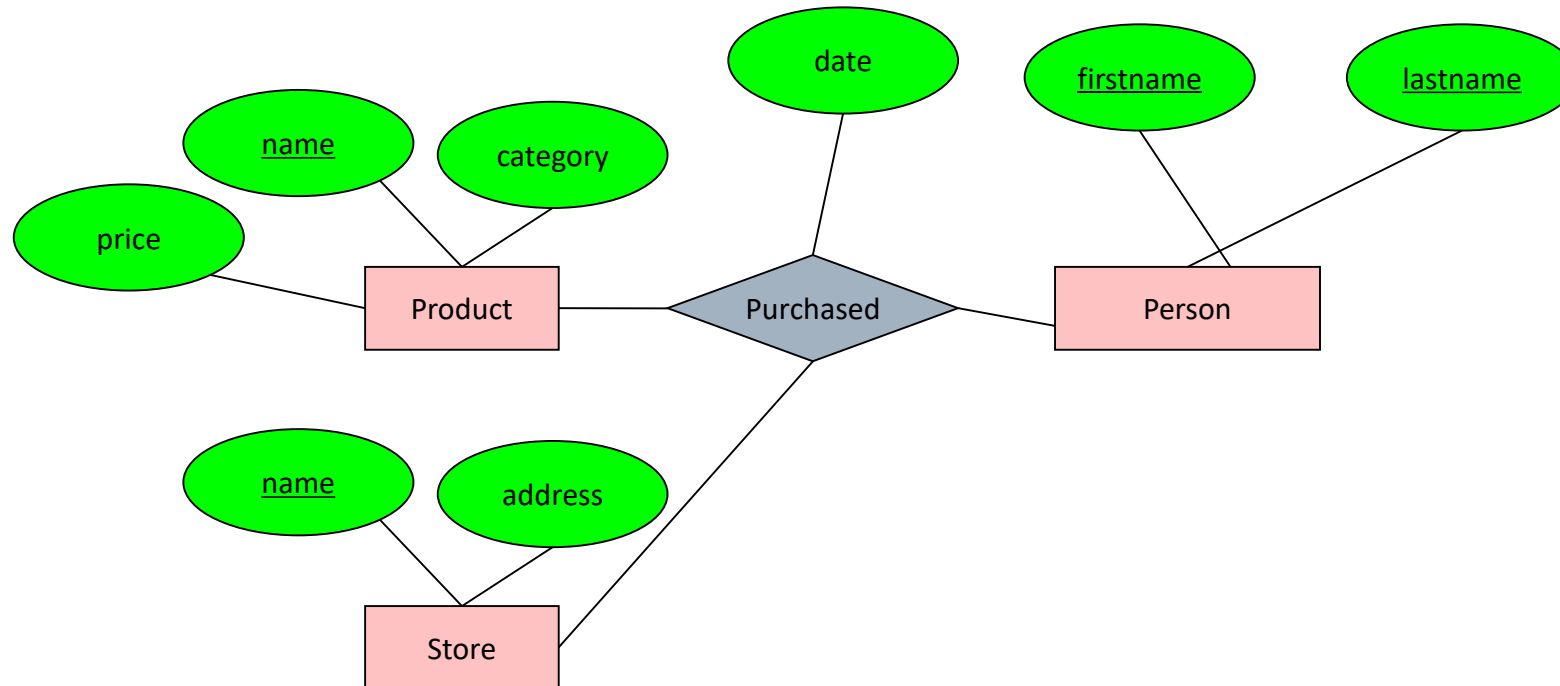
```
CREATE TABLE Purchased(  
  name CHAR(50),  
  firstname CHAR(50),  
  lastname CHAR(50),  
  date DATE,  
  PRIMARY KEY (name, firstname, lastname),  
  FOREIGN KEY (name)  
    REFERENCES Product,  
  FOREIGN KEY (firstname, lastname)  
    REFERENCES Person  
)
```



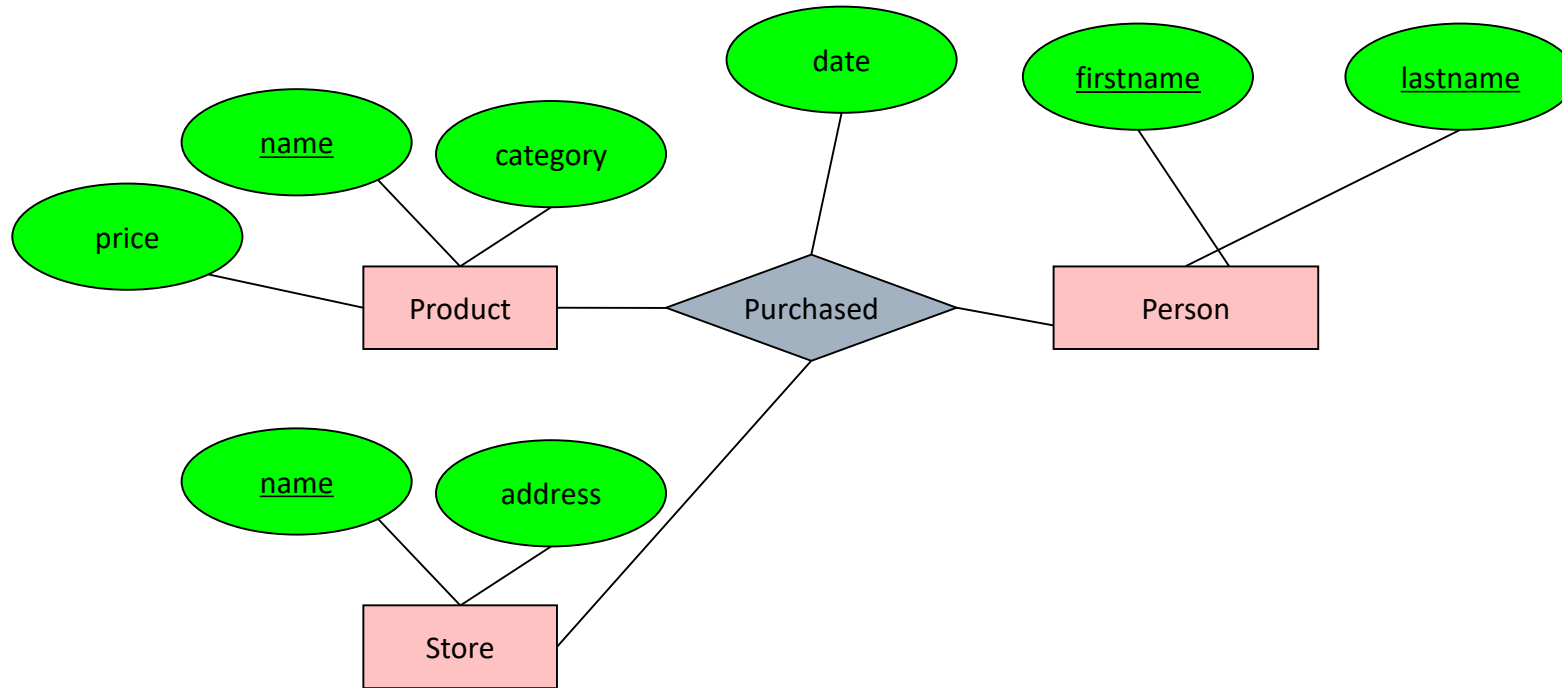
Purchased			
<u>name</u>	<u>firstname</u>	<u>lastname</u>	date
Gizmo1	Bob	Alice	01/01/15
Gizmo2	Alice	Bob	01/03/15
Gizmo1	Joe	Smith	01/05/15

From E/R Diagram to Relational Schema

How do we represent this as a relational schema?



From E/R Diagram to Relational Schema



Product

<u>Name</u>	Price	Category
Gizmo1	99.99	Camera
Gizmo2	19.99	Edible

Purchased

<u>Pname</u>	<u>Firstname</u>	<u>Lastname</u>	Date
Gizmo1	Bob	Alice	01/01/15
Gizmo2	Alice	Bob	01/03/15
Gizmo1	Joe	Smith	01/05/15



Practice 2

Add arrows to your E/R diagram!

Also make sure to add (new concepts underlined):



A player can only belong to one team, a play can only be in one game, a pass/run..?



Players can achieve a **Personal Record** linked to a specific Game and Play



Players have a **weight** which changes in on vs. off-season

Lab Assignment #1

- Released date: 3/16
- Due date: 3/29 23:59 PM
- Where to submit: to e-class (<http://eclass.seoultech.ac.kr>)
- Submission: a word file consisting of 1) the result ER diagram and 2) its explanation
- Late submission is not allowed
- Description
 - Practice 2 and Practice 3
 - Based on the result of Practice 1 explained in the lecture, extend the relationships, entity sets, attributes according to the requirements in Practices 2 and 3
 - There could be many possible answers. The important check point will be if the results are made by yourself.

Today's Lecture


1. E/R Basics: Entities & Relations
2. E/R Design considerations
3. **Advanced E/R Concepts**

What you will learn about in this section

- 1. Subclasses**
- 2. Constraints**
- 3. Weak entity sets**

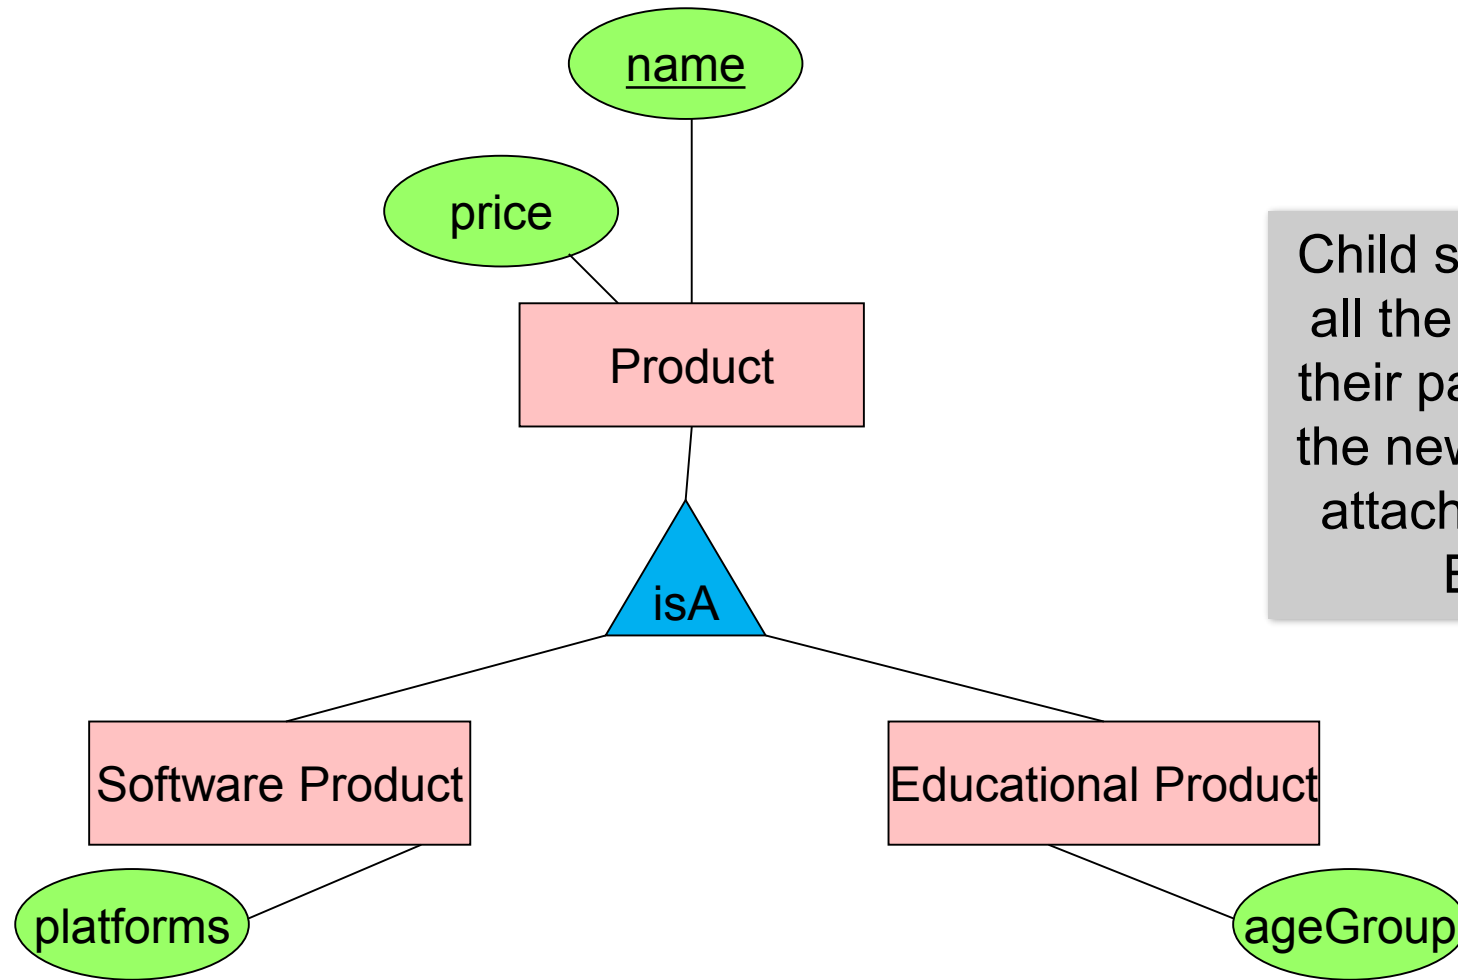
Modeling Subclasses

- Some objects in a class may be special, i.e. worthy of their own class
- Define a new class?
 - *But what if we want to maintain connection to current class?*
- Better: define a subclass
 - *Ex:*



We can define **subclasses** in E/R!

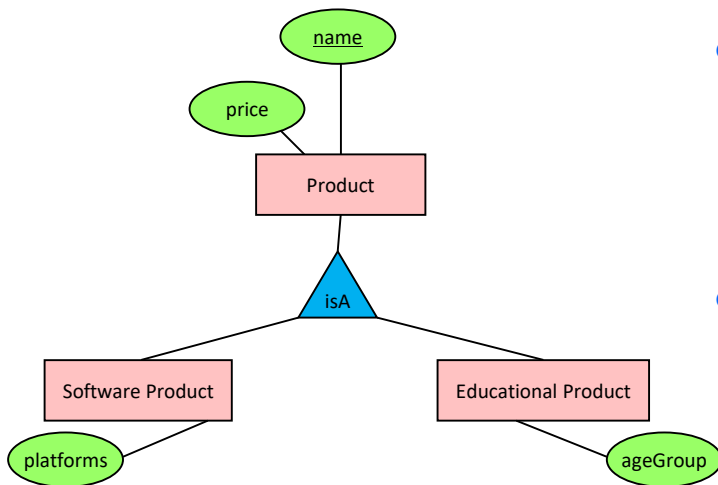
Modeling Subclasses



Child subclasses contain all the attributes of *all* of their parent classes **plus** the new attributes shown attached to them in the E/R diagram

Understanding Subclasses

■ Think in terms of records; ex:



- Product
- SoftwareProduct
- EducationalProduct

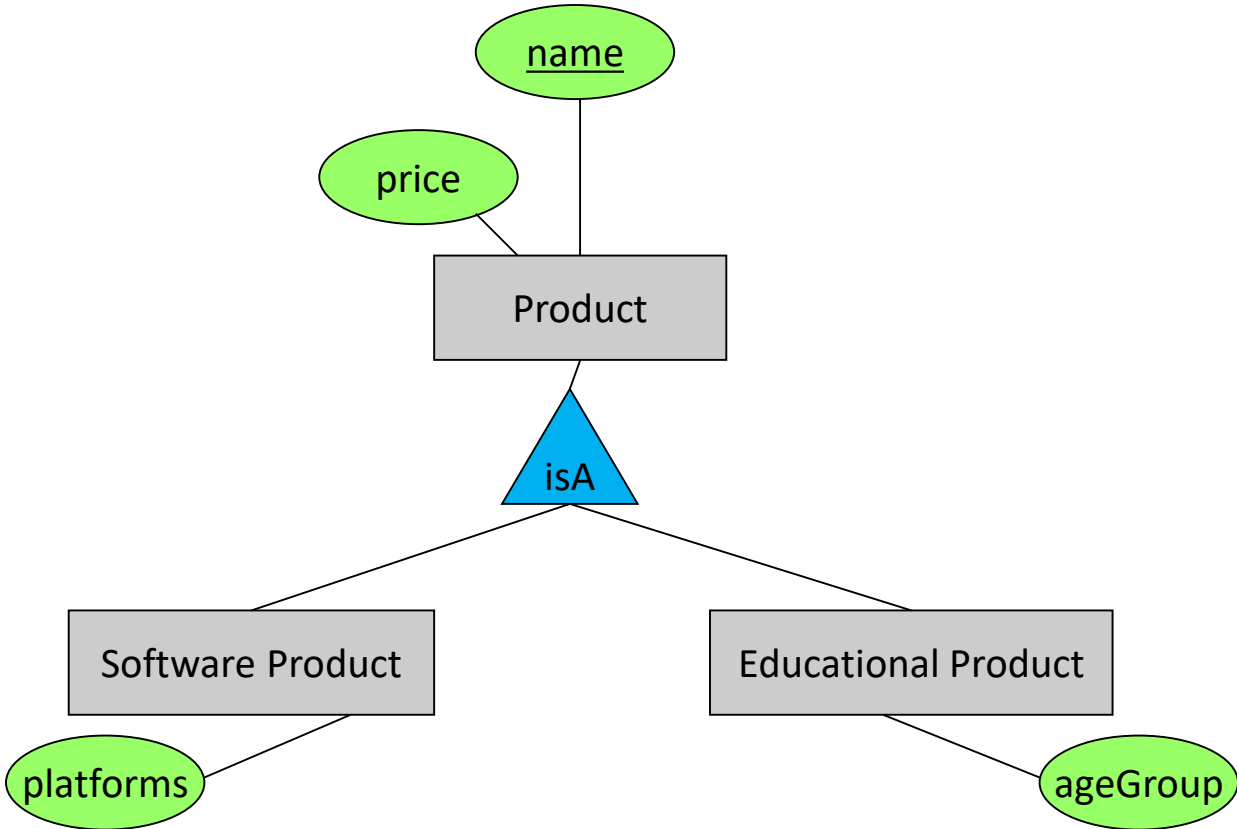
name
price

name
price
platforms

name
price
ageGroup

Child subclasses contain all the attributes of *all* of their parent classes **plus** the new attributes shown attached to them in the E/R diagram

Think like tables...



Product

<u>name</u>	price	category
Gizmo	99	gadget
Camera	49	photo
Toy	39	gadget

Sw.Product

<u>name</u>	platforms
Gizmo	unix


Ed.Product

<u>name</u>	ageGroup
Gizmo	toddler
Toy	retired

IsA Review

- If we declare *A IsA B* then every A is a B
- We use IsA to Add descriptive attributes to a subclass

Modeling UnionTypes With Subclasses




Suppose each piece of furniture is owned either by a person, or by a company. *How do we represent this?*

Modeling Union Types with Subclasses

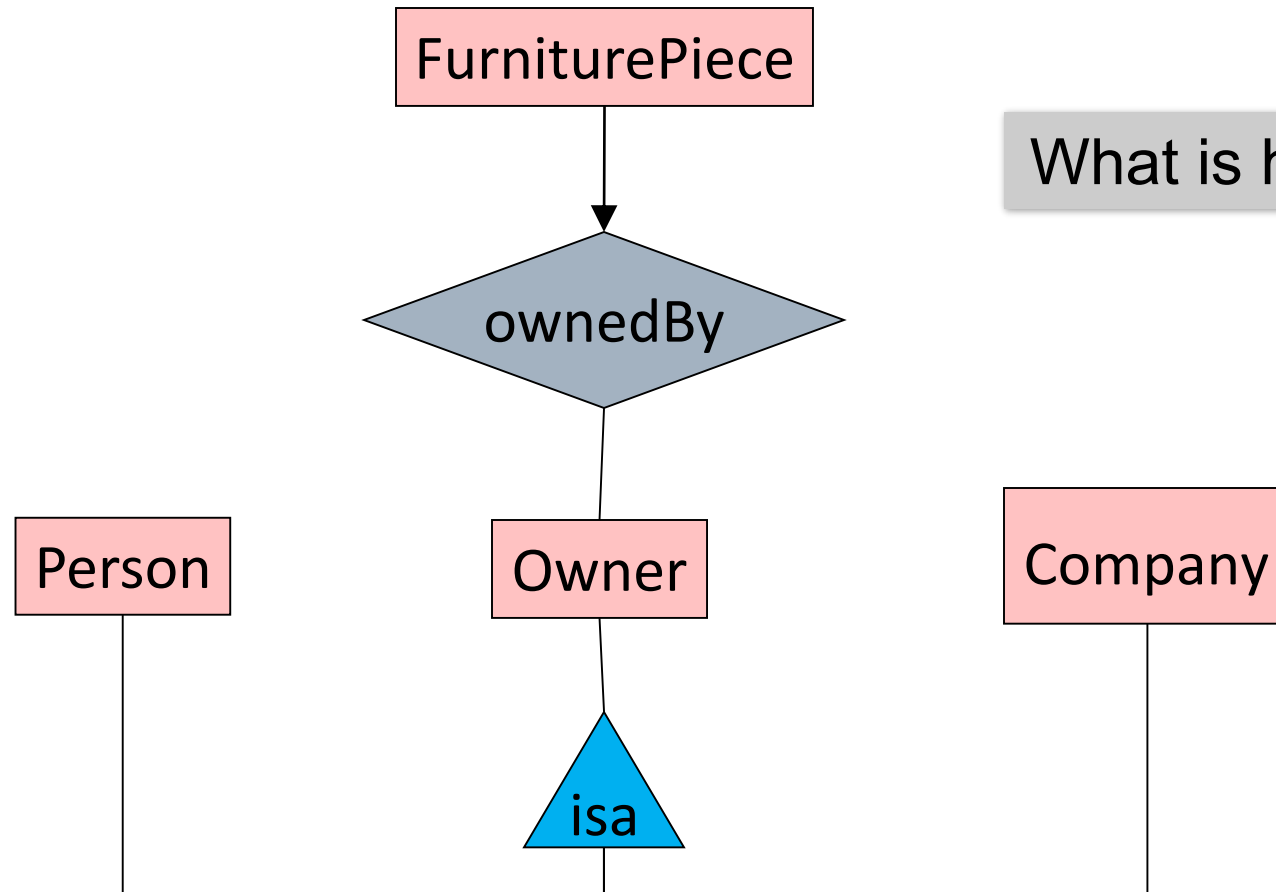
Say: each piece of furniture is owned either by a person, or by a company

Solution 1. Acceptable, but imperfect (What's wrong ?)



Modeling Union Types with Subclasses

Solution 2: better (though more laborious)




What is happening here?

Constraints in E/R Diagrams


- Finding constraints is part of the E/R modeling process. Commonly used constraints are:
 - Keys: Implicit constraints on uniqueness of entities
 - *Ex: An SSN uniquely identifies a person*
 - Single-value constraints:
 - *Ex: a person can have only one father*
 - Referential integrity constraints: Referenced entities must exist
 - *Ex: if you work for a company, it must exist in the database*
 - Other constraints:
 - *Ex: peoples' ages are between 0 and 150*

Recall
FOREIGN
KEYs!

Participation Constraints: Partial v. Total



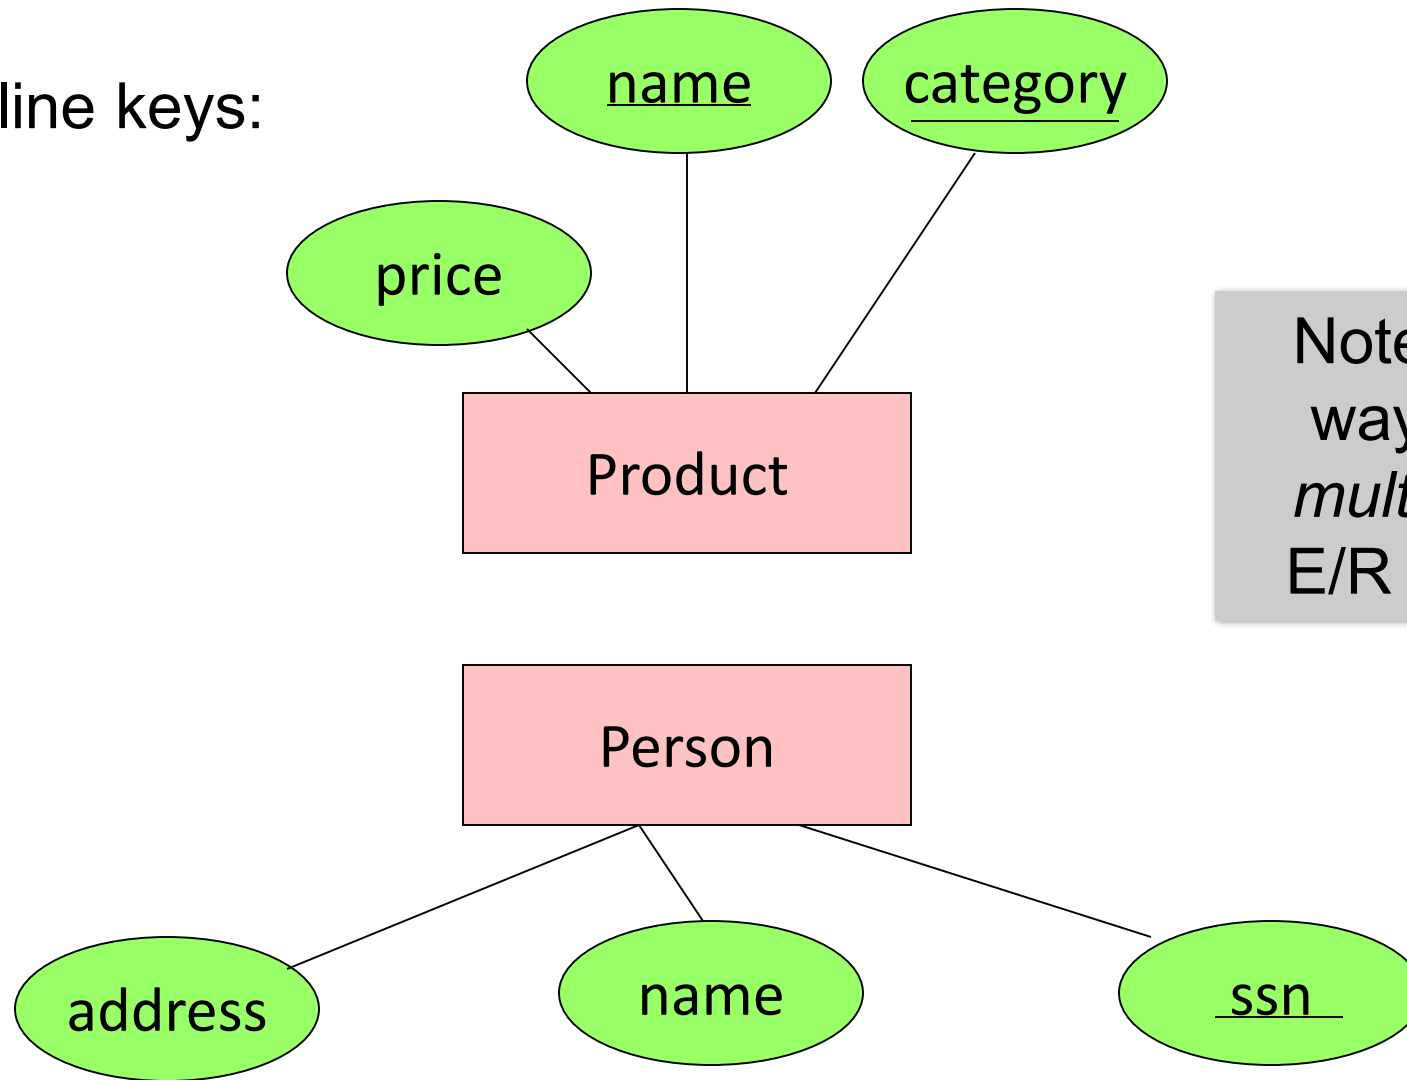
Are there products made by no company?
Companies that don't make a product?



Bold line indicates total participation (i.e. here: all products are made by a company)


Keys in E/R Diagrams

Underline keys:




Note: no formal way to specify *multiple* keys in E/R diagrams...


Single Value Constraints




v. s.



Referential Integrity Constraints




Each product made by at most one company.
Some products made by no company?



Each product made by exactly one company.

Weak Entity Sets


Entity sets are weak when their key comes from other classes to which they are related.



“Football team” v. “***The Stanford*** Football team” (E.g., *Berkeley has a football team too, sort of*)

Weak Entity Sets

Entity sets are weak when their key comes from other classes to which they are related.




- number is a partial key. (denote with dashed underline).
- University is called the identifying owner.
- Participation in affiliation must be total. Why?



Practice 3

Weak entity sets / Subclasses

Concepts to include / model:



Teams belong
to cities- model
as ***weak entity
sets***



Players are either
on Offense or
Defense, and are
of types (QB, RB,
WR, TE, K,
Farmer*...)

E/R Summary

- **E/R diagrams are a visual syntax that allows technical and non-technical people to talk**
 - For conceptual design
- **Basic constructs: entity, relationship, and attributes**
- **A good design is faithful to the constraints of the application**