



Transactions

Prof. Hyuk-Yoon Kwon

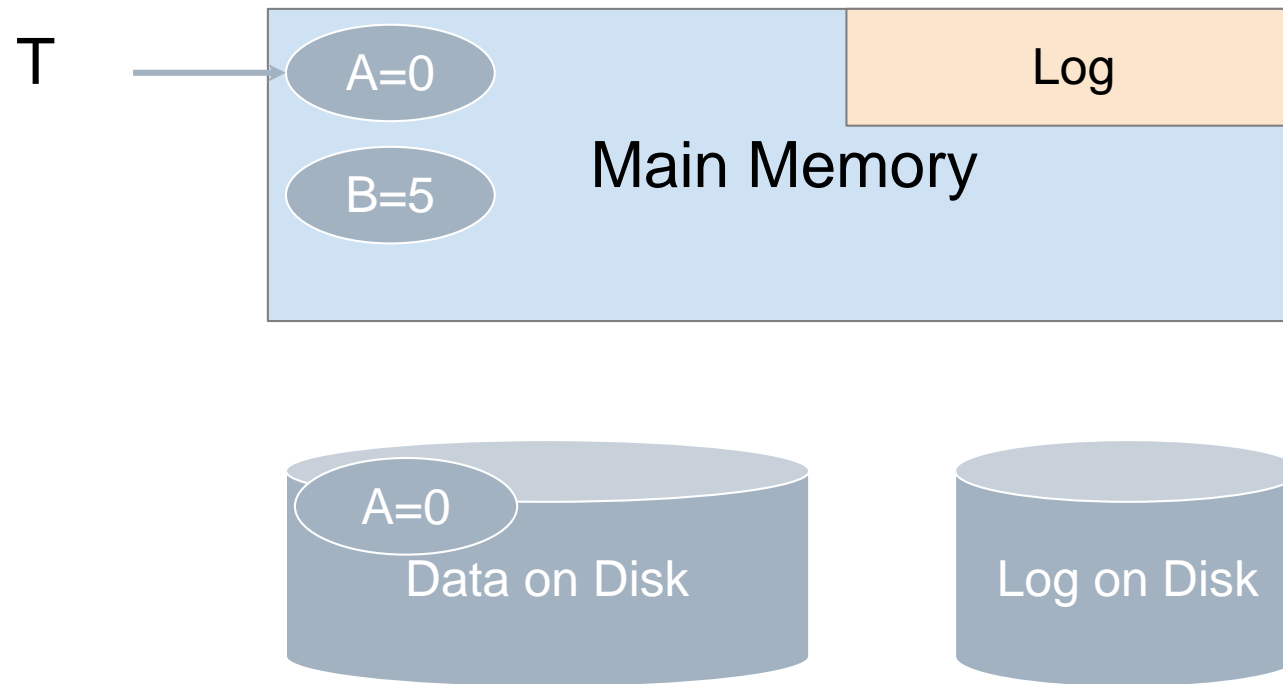
<https://sites.google.com/view/seoultech-bigdata>

Most parts are based on slides used in Stanford (<http://web.stanford.edu/class/cs145>)

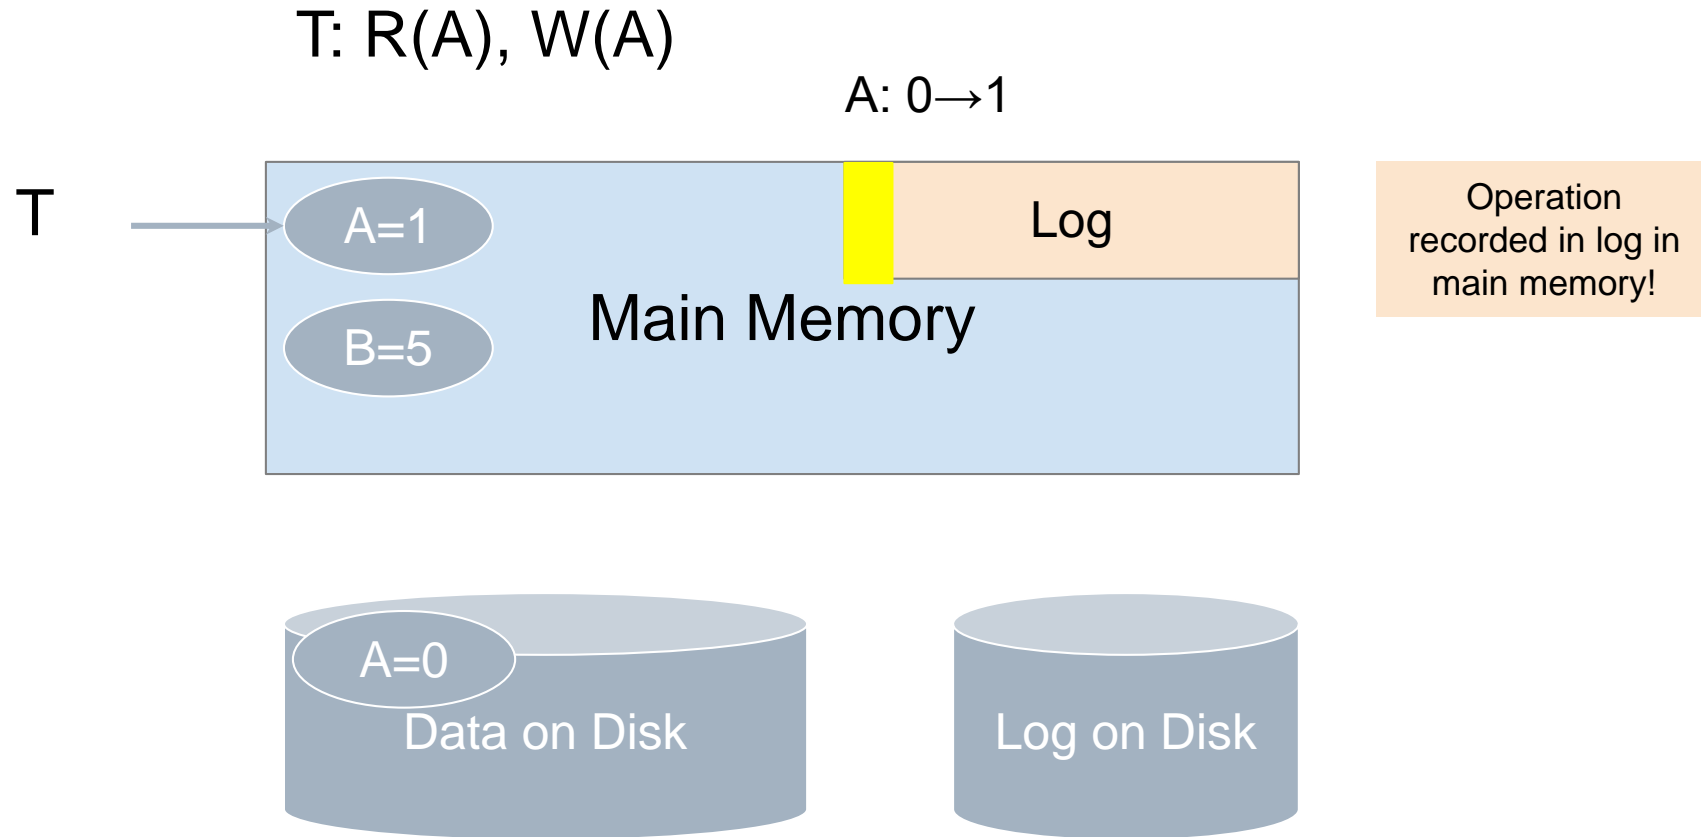
Atomicity & Durability via Logging

A picture of logging

T: R(A), W(A)



A picture of logging



What is the correct way to write this all to disk?

- We'll look at the *Write-Ahead Logging (WAL)* protocol
- We'll see why it works by looking at other protocols which are incorrect!

Remember: Key idea is to ensure durability
while maintaining our ability to “undo”!

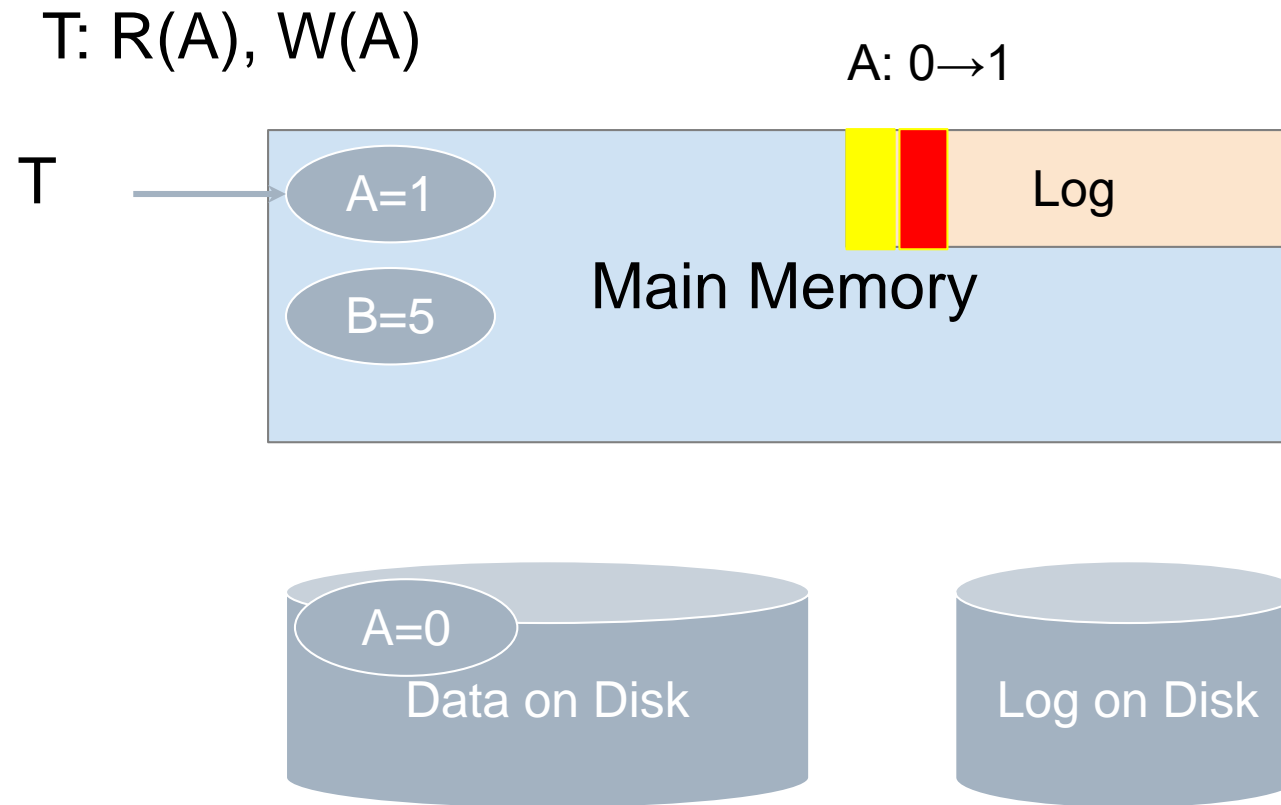
Write-Ahead Logging (WAL) TXN Commit Protocol

Transaction Commit Process

- FORCE Write **commit** record to log
- All log records up to last update from this TX are FORCED
- Commit() returns

Transaction is committed *once commit log record is on stable storage*

Incorrect Commit Protocol #1



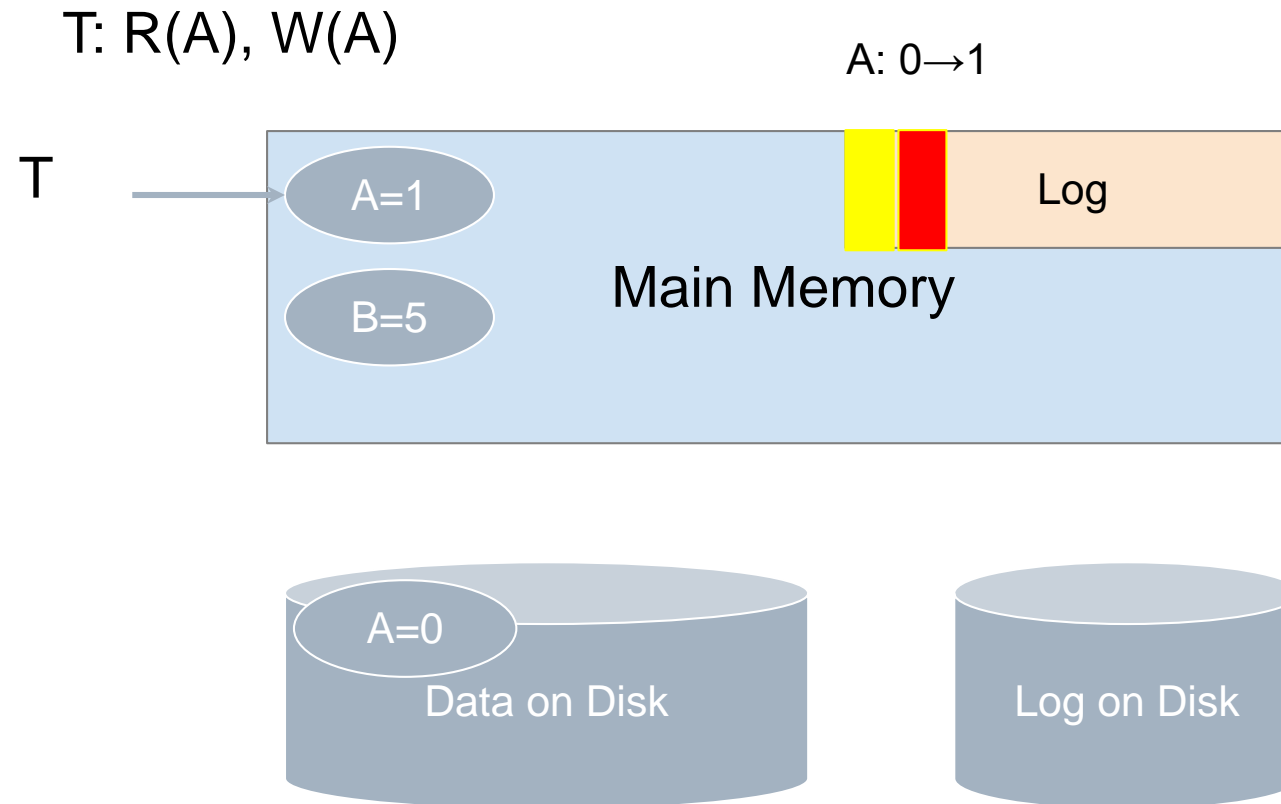
Let's try committing
before we've written
either data or log to
disk...

OK, Commit!

If we crash now, is T
durable?

Lost T's update!

Incorrect Commit Protocol #2



Let's try committing *after* we've written data but *before* we've written log to disk...

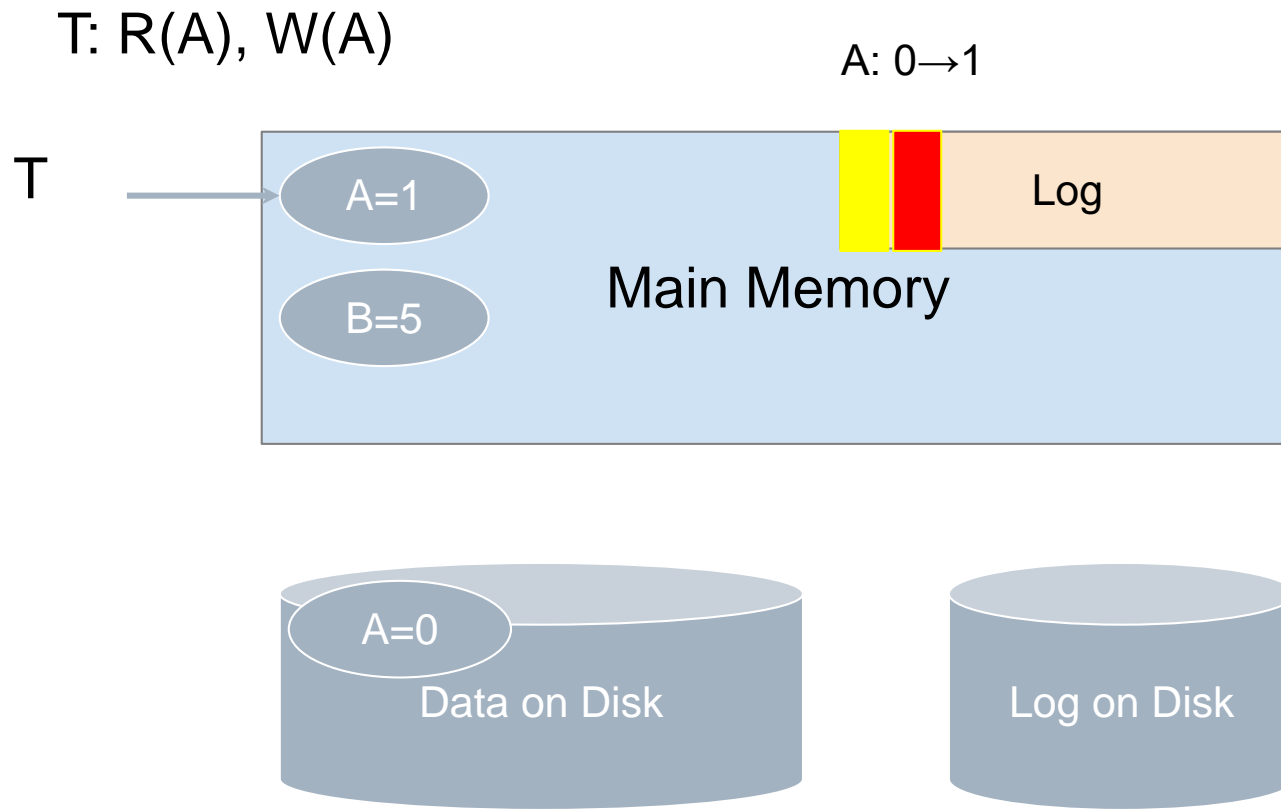
OK, Commit!

If we crash now, is T durable? Yes! Except...

How do we know whether T was committed??

Improved Commit Protocol (WAL)

Write-ahead Logging (WAL) Commit Protocol

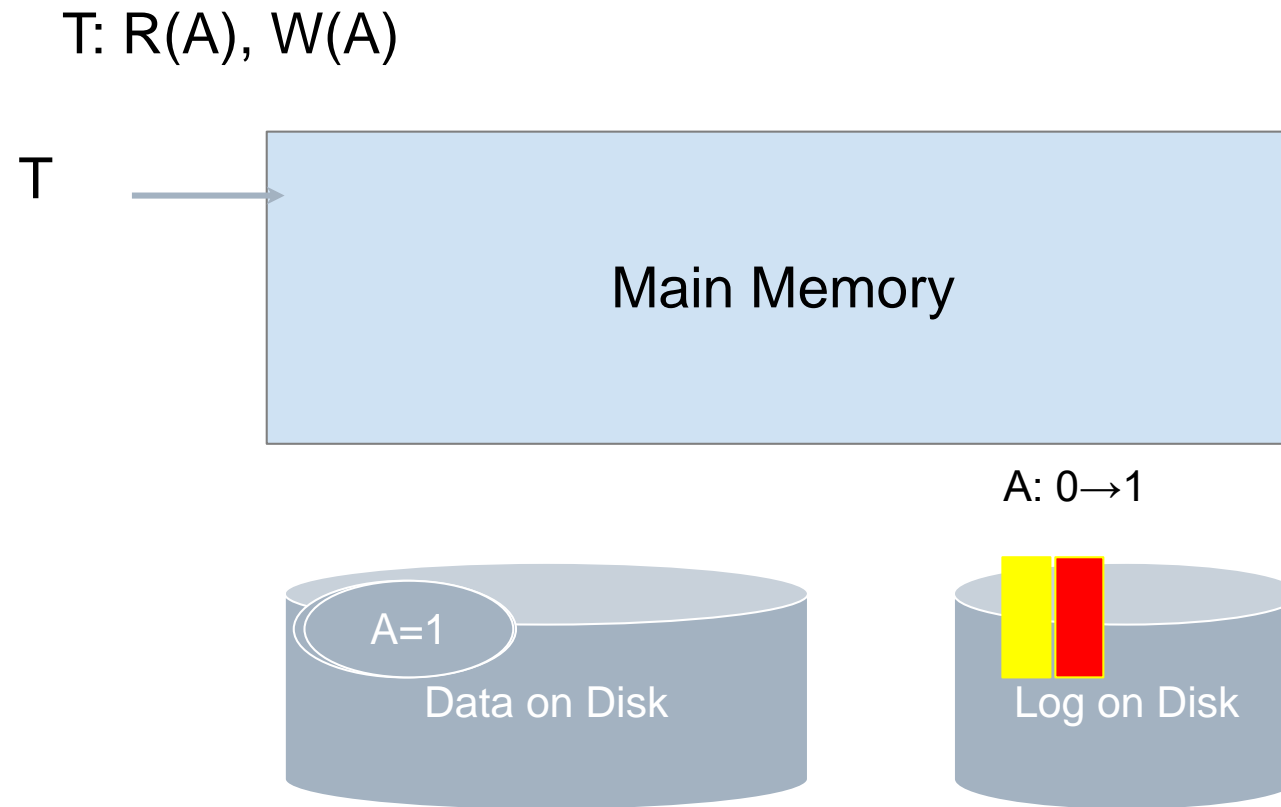


This time, let's try committing after we've written log to disk but before we've written data to disk... this is WAL!

OK, Commit!

If we crash now, is T durable?

Write-ahead Logging (WAL) Commit Protocol



This time, let's try committing after we've written log to disk but before we've written data to disk... this is WAL!

OK, Commit!

If we crash now, is T durable?

USE THE LOG!

Write-Ahead Logging (WAL)

Algorithm: WAL

1. Must *force log record* for an update *before* the corresponding data page goes to storage
2. Must write all log records for a TX before commit

→ **Atomicity**

→ **Durability**

Logging Summary

- If DB says TX **commits**, TX effect **remains** after database crash
- DB can **undo actions** and help us with **atomicity**

HW Assignment #1

■ Directions

1. We will now create a new table named "T", having three columns: id (of type integer, the primary key), s (of type character string with a length varying from 1 to 40 characters), and si (of type small integer)
2. Then, insert some rows to the newly created table:
 - INSERT INTO T (id, s) VALUES (1, 'first');
 - INSERT INTO T (id, s) VALUES (2, 'second');
 - INSERT INTO T (id, s) VALUES (3, 'third');
3. Check the data stored in T
4. Try to cancel or rollback the current transaction and check the data in table T
 - **Capture the screen showing the results**
5. Try to execute "commit;" command before rollback the transaction in 4 and compare the result with it in 4
 - **Capture the screen showing the results**

■ Submit two screen shots above

Example

Monthly bank
interest
transaction

Full run

Money

Account	Balance (\$)
3001		500
4001		100
5001		20
6001		60
3002		80
4002		-200
5002		320
...		...
30108		-100
40008		100
50002		20

Money (@4:29 am day+1)

Account	Balance (\$)
3001		550
4001		110
5001		22
6001		66
3002		88
4002		-220
5002		352
...		...
30108		-110
40008		110
50002		22

WAL (@4:29 am day+1)

T-Monthly-423	START TRANSACTION		
T-Monthly-423	3001	500	550
T-Monthly-423	4001	100	110
T-Monthly-423	5001	20	22
T-Monthly-423	6001	60	66
T-Monthly-423	3002	80	88
T-Monthly-423	4002	-200	-220
T-Monthly-423	5002	320	352
T-Monthly-423
T-Monthly-423	30108	-100	-110
T-Monthly-423	40008	100	110
T-Monthly-423	50002	20	22
T-Monthly-423	COMMIT		

'T-Monthly-423'

Monthly Interest 10%

4:28 am Starts run on 10M bank accounts

Takes 24 hours to run

START TRANSACTION

UPDATE Money

SET Amt = Amt * 1.10

COMMIT

Example

Monthly bank

interest

transaction

With crash

Money

Account	Balance (\$)
3001		500
4001		100
5001		20
6001		60
3002		80
4002		-200
5002		320
...		...
30108		-100
40008		100
50002		20

Money (@10:45 am)

Account	Balance (\$)
3001		550
4001		110
5001		22
6001		66
3002		88
4002		-200
5002		320
...		...
30108		-110
40008		110
50002		22

??

??

??

??

WAL log (@10:29 am)

T-Monthly-423	START TRANSACTION		
T-Monthly-423	3001	500	550
T-Monthly-423	4001	100	110
T-Monthly-423	5001	20	22
T-Monthly-423	6001	60	66
T-Monthly-423	3002	80	88
T-Monthly-423
T-Monthly-423	30108	-100	-110
T-Monthly-423	40008	100	110
T-Monthly-423	50002	20	22
T-Monthly-423	4002	-200	-220
T-Monthly-423	5002	320	352

'T-Monthly-423'

Monthly Interest 10%

4:28 am Starts run on 10M bank accounts

Takes 24 hours to run

Network outage at 10:29 am,

System access at 10:45 am

Did T-Monthly-423 complete?

Which tuples are bad?

Case1: T-Monthly-423 was crashed

Case2: T-Monthly-423 completed. 4002 deposited 20\$ at 10:45 am

Example

Monthly bank interest transaction

Recovery

Money (@10:45 am)

Account	Balance (\$)
3001		550
4001		110
5001		22
6001		66
3002		88
4002		-200
5002		320
...		
30108		-110
40008		110
50002		22

Money (after recovery)

Account	Balance (\$)
3001		500
4001		100
5001		20
6001		60
3002		80
4002		-200
5002		320
...		...
30108		-100
40008		100
50002		20

WAL log (@10:29 am)

T-Monthly-423	START TRANSACTION		
T-Monthly-423	3001	500	550
T-Monthly-423	4001	100	110
T-Monthly-423	5001	20	22
T-Monthly-423	6001	60	66
T-Monthly-423	3002	80	88
T-Monthly-423
T-Monthly-423	30108	-100	-110
T-Monthly-423	40008	100	110
T-Monthly-423	50002	20	22
T-Monthly-423	4002	-200	-220
T-Monthly-423	5002	320	352

System recovery (after 10:45 am)

- 1 Rollback uncommitted transactions
 - Restore old values from WALlog (if any)
 - Notify developers about aborted txn
- 1.1 Redo Recent transactions (w/ new values)
- 2 Redo (any pending) transactions

Example

Monthly bank

interest

transaction

Performance

Money

Account	Balance (\$)
3001		500
4001		100
5001		20
6001		60
3002		80
4002		-200
5002		320
...		...
30108		-100
40008		100
50002		20

Money (@4:29 am day+1)

Account	Balance (\$)
3001		550
4001		110
5001		22
6001		66
3002		88
4002		-220
5002		352
...		...
30108		-110
40008		110
50002		22

WAL (@4:29 am day+1)

T-Monthly-423	START TRANSACTION		
T-Monthly-423	3001	500	550
T-Monthly-423	4001	100	110
T-Monthly-423	5001	20	22
T-Monthly-423	6001	60	66
T-Monthly-423	3002	80	88
T-Monthly-423	4002	-200	-220
T-Monthly-423	5002	320	352
T-Monthly-423
T-Monthly-423	30108	-100	-110
T-Monthly-423	40008	100	110
T-Monthly-423	50002	20	22
T-Monthly-423	COMMIT		

Cost to update all data

10M bank accounts → 10M seeks? (worst case)

(@10 msec/seek, that's 100,000 secs)

Speedup for commit
100,000 secs vs 1 sec!!!



Cost to Append to log

+ 1 seek to get 'end of log'


+ write 10M log entries sequentially (fast!)

(< 1 sec)

[Lazily update data on disk later, when convenient.]

Our primary model [recap]

- **Shared:** Each process can read from / write to shared data in main memory





- **Disk:** Global memory can read from / flush to disk
- **Log:** Assume on stable disk storage- spans both main memory and disk...

Log is a *sequence* from main memory -> disk

“Flushing to disk” = writing to disk from main memory

Cluster model

- A popular alternative (with tradeoffs)



Commit by replicating log and/or data
to 'n' other machines (e.g. $n = 2$)

[different rack or different datacenter]

Example

Failure model

Main model: RAM could fail, Disk is durable

VS

Cluster model: “Architect” your hardware to be fault resilient
RAM on different machines don’t fail at same time
Power to racks is uncorrelated

Incremental cost to write to machine

Network speeds intra-datacenter could be 1-10 microsecs

[Lazily update data on disk later, when convenient]



Example: Youtube DB

YouTube

funny cats

Home

Trending

Subscriptions

Library

History

Watch later

Purchases7

Liked videos

Subscriptions

Popular on YouTube...

Music

Sports

Gaming

From YouTube

YouTube Premium

Movies & Shows

About 12,100,000 results

REMOVE CAT PEE STAINS

1:46

How to Get Rid of Cat Pee Stains

Ad BISSELL • 2M views

Your cat had an accident on the carpet. BISSELL is here to help!

CATS make us LAUGH ALL THE TIME! - Ultra FUNNY CAT

Tiger FunnyWorks • 100K views • 3 days ago

Ultra funny cats and kitten that will make you cry with laughter! Cats are the best laugh all the time! This is ...

New

You will LAUGH SO HARD that YOU WILL FAINT - FUNNY C compilation

Tiger FunnyWorks • 19M views • 8 months ago

Well well well, cats for you again. But this time, even better, even funnier, even more you like these furries the ...

Have you EVER LAUGHED HARDER? - Ultra FUNNY CATS

Tiger FunnyWorks • 124K views • 1 week ago

Super funny cats and kitten that will make you scream with laughter! This is the ! LAUGH challenge ever!

cats funny

Upload video

Go live

Shop Now >> DressLily

Up next

Top Cats Vs. Cucumbers

Funny Cat Videos Compilation.

Animal Planet Videos • 23M views

Funny Elias play with the wheel on the bus and another toys - ...

Elias Adventures

291 watching

LIVE NOW

LIVE: Rescue kitten nursery! TinyKittens.com

TinyKittens HQ

1.3K watching

LIVE NOW

Puss in boots and the three diablos [HD]

Mathew Garcia

7.6M views

Animal Planet Videos

Published on May 23, 2018

Baby Cats - Funny and Cute Baby Cat Videos Compilation (2018) Gatitos Bebes Video Recopilación

Animal Planet Videos

Subscribe Here: <https://goo.gl/qor4XN>

SHOW MORE

809,337 views

4.7K

768


SHARE

SUBSCRIBE 337K

Example

Design 1: WAL Log for Video Views
<videoid, old # views, new # views>

T-LIKE-4307	START TRANSACTION		
T-LIKE-4307	3001	537	538
T-LIKE-4307	COMMIT		
T-LIKE-4308	START TRANSACTION		
T-LIKE-4308	5309	100001	10002
T-LIKE-4308	COMMIT		
T-LIKE-4309	START TRANSACTION		
T-LIKE-4309	3001	538	539
T-LIKE-4309	COMMIT		
...
T-LIKE-4341	5309	100002	10003
T-LIKE-4351	5309	100003	10004
...
T-LIKE-4383	START TRANSACTION		
T-LIKE-4383	5309	100004	10005
T-LIKE-4383	COMMIT		



Luis Fonsi - Despacito ft. Daddy Yankee

5,611,744,868 views


30M 3.5M

SHARE SAVE ...

Example

Design 2: Replicate #Video Views in cluster
<unix time, videoid, # views>


1539893189	3001	'+1'
1539893195	5309	'+1'
1539893225	3001	'+1'
..		'+1'
	5309	'+1'
...		'+1'
	5309	'+1'
...		'+1'
1539893289	5309	'+1'



Write to RAM on
n=3 machines
(<videoid, #likes>)



Example

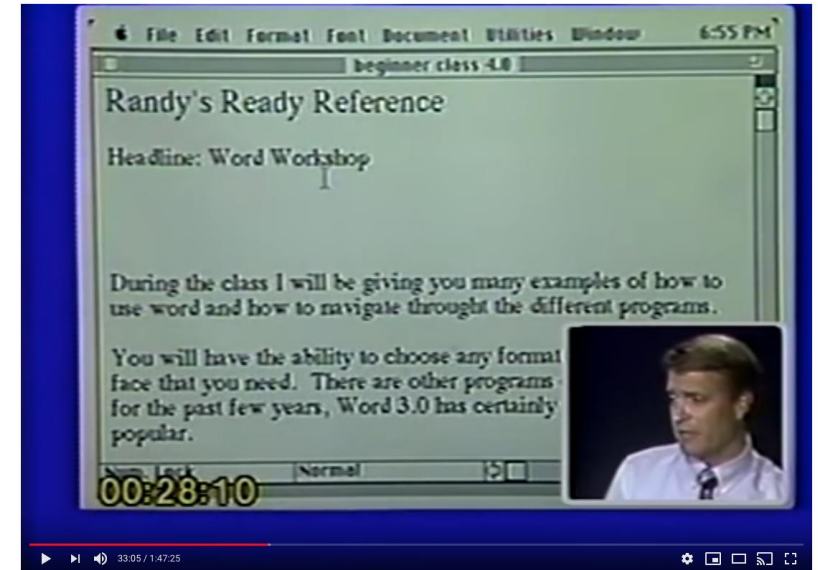


Popular video

Design #3

For most videos, Design 1 (full WAL logs)

For popular videos, Design 2



Unpopular video

Summary

Design Questions?

Correctness: Need true ACID? Pseudo-ACID? What losses are OK?

Design parameters:

Any data properties you can exploit? (e.g., '+1', popular vs not)

How much RAM, disks and machines?

How many writes per sec?

How fast do you want system to recover?

Choose: WAL logs, Replication on n-machines, Hybrid?

Concurrency & Locking

What you will learn about in this section

- Interleaving & scheduling

- Conflict & anomaly types

Note: Go back to our simple single machine model for RAM/disk for now

Concurrency: Isolation & Consistency

- The DBMS must handle concurrency such that...

- **Isolation** is maintained: Users must be able to execute each TXN **as if they were the only user**
 - DBMS handles the details of *interleaving* various TXNs

ACID

- **Consistency** is maintained: TXNs must leave the DB in a **consistent state**
 - DBMS handles the details of enforcing integrity constraints

ACCID

Note the hard part...

...is the effect of *interleaving* transactions and *crashes*.

Example- consider two TXNs:

T1: START TRANSACTION

UPDATE Accounts
SET Amt = Amt + 100
WHERE Name = 'A'

UPDATE Accounts
SET Amt = Amt - 100
WHERE Name = 'B'

COMMIT

T1 transfers \$100 from B's account to A's account

T2: START TRANSACTION

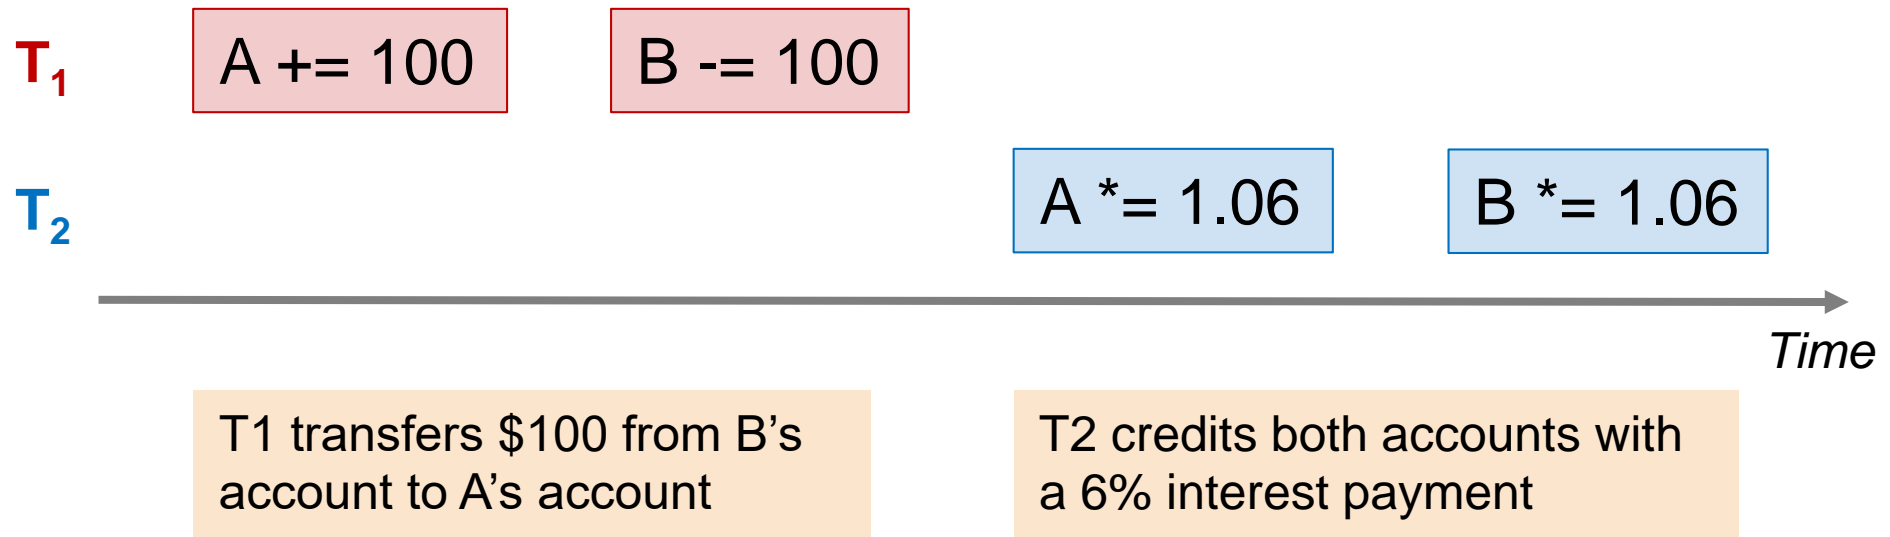
UPDATE Accounts
SET Amt = Amt * 1.06

COMMIT

T2 credits both accounts with a 6% interest payment

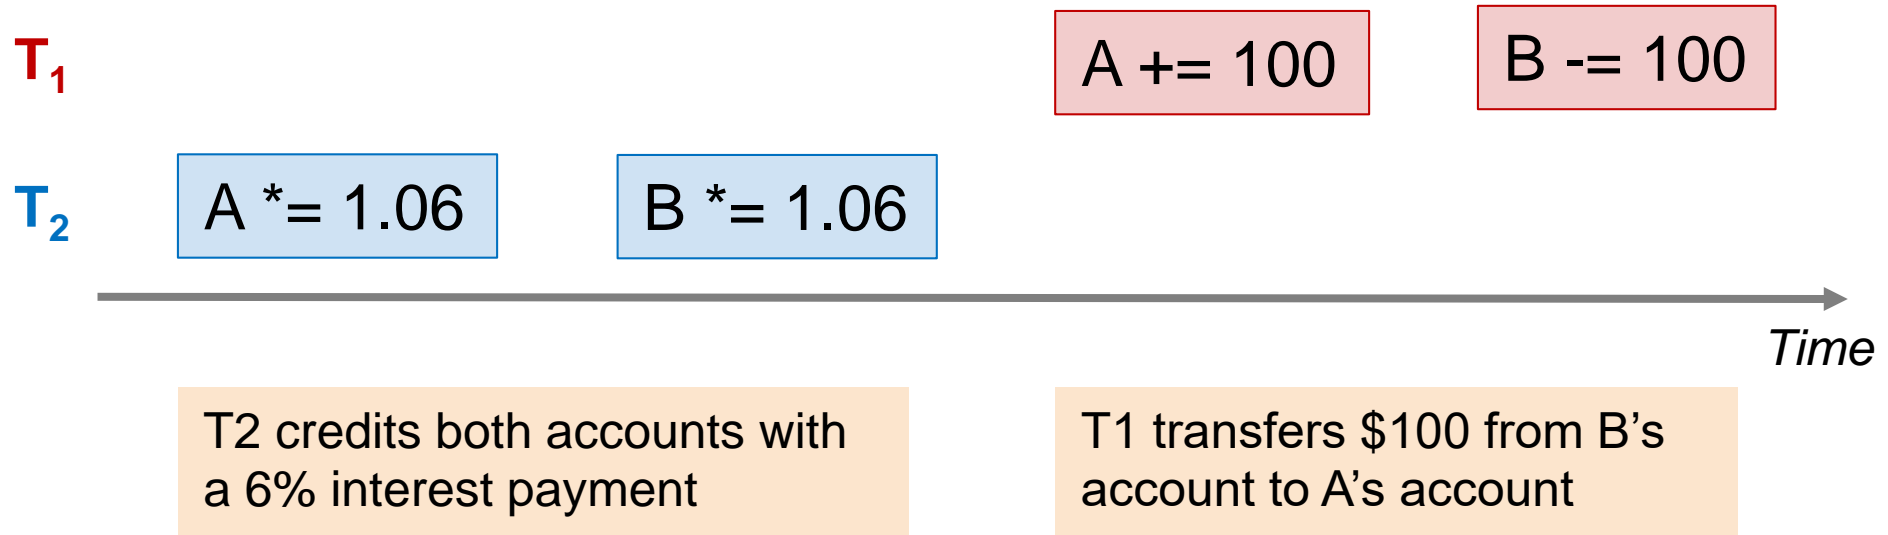
Example- consider two TXNs:

We can look at the TXNs in a timeline view- serial execution:



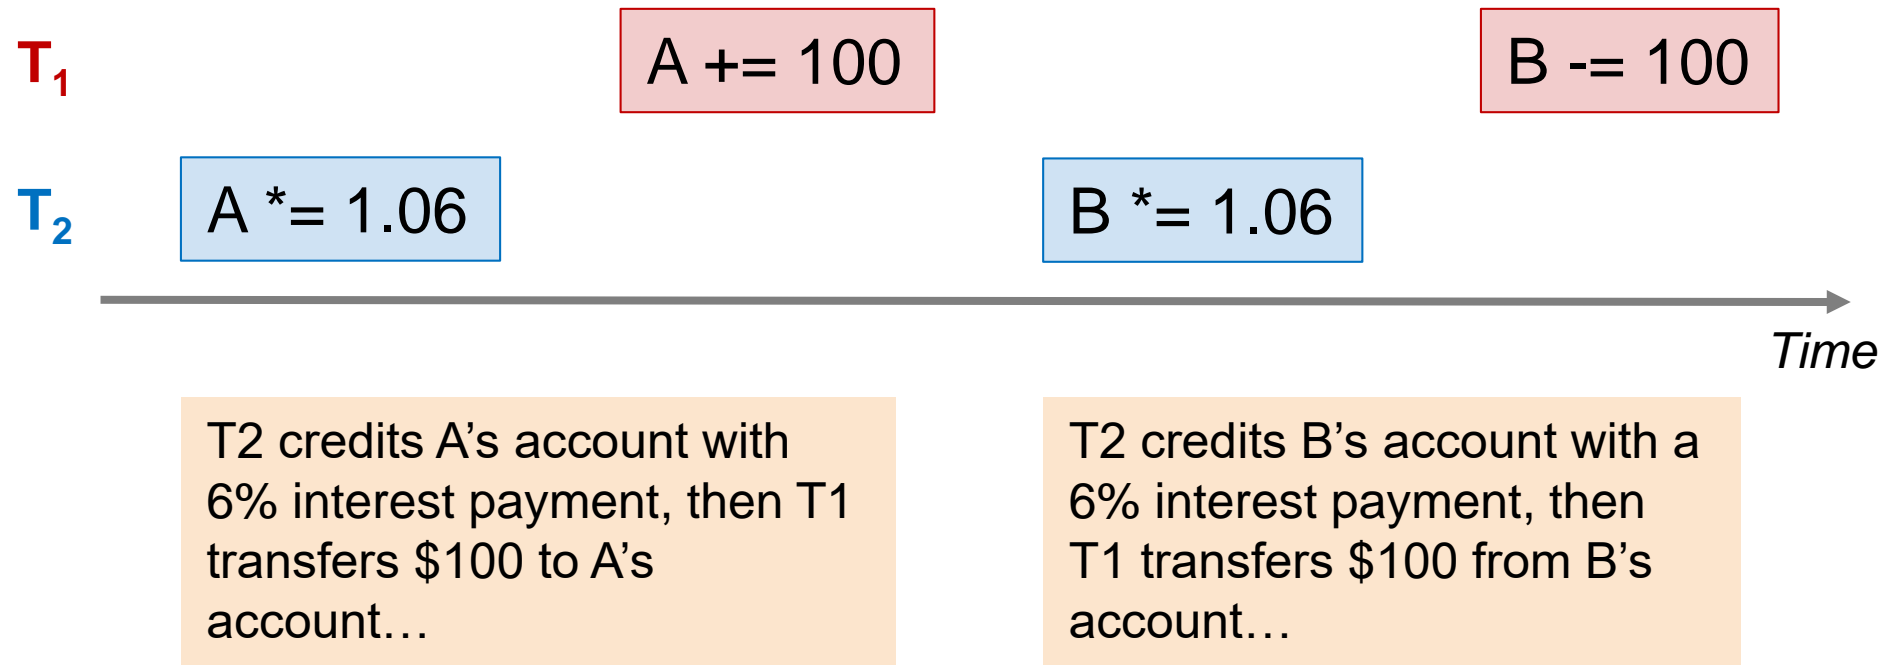
Example- consider two TXNs:

The TXNs could occur in either order... DBMS allows!



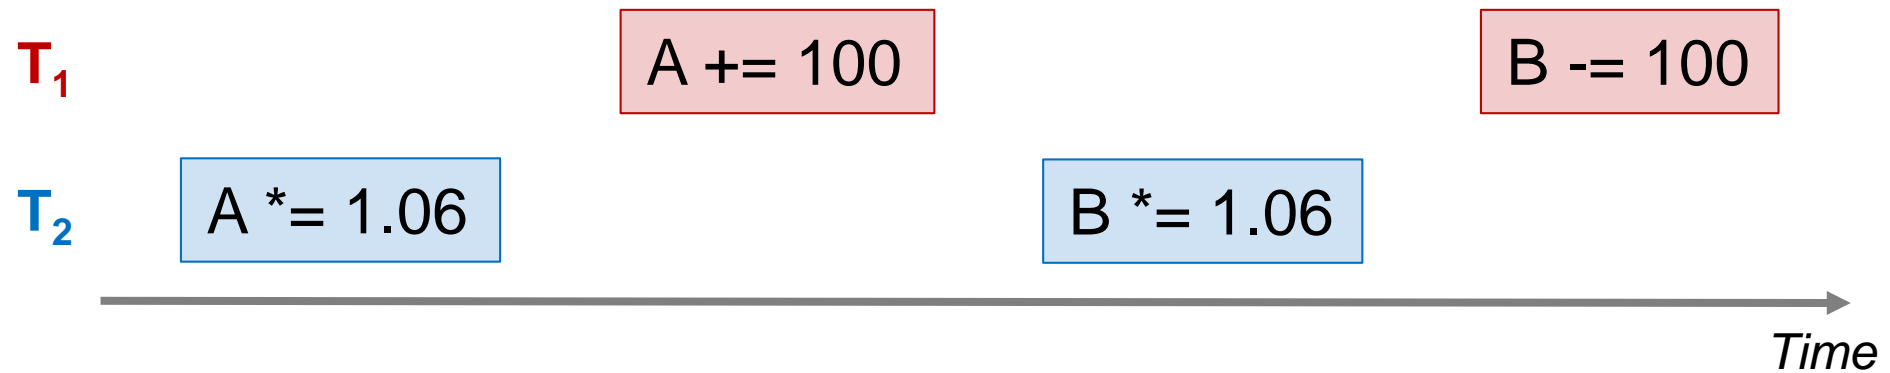
Example- consider two TXNs:

The DBMS can also **interleave** the TXNs



Example- consider two TXNs:

The DBMS can also **interleave** the TXNs



What goes wrong here??

Why Interleave TXNs?

- Interleaving TXNs might lead to anomalous outcomes... why do it?
- Several important reasons:
 - Individual TXNs might be *slow*- don't want to block other users during!
 - Disk access may be *slow*- let some TXNs use CPUs while others accessing disk!

All concern large differences in ***performance***

Interleaving & Isolation

- The DBMS has freedom to interleave TXNs
- However, it must pick an interleaving or **schedule** such that isolation and consistency are maintained
- Must be *as if* the TXNs had executed serially!

DBMS must pick a schedule which maintains isolation & consistency

Scheduling examples

*Starting
Balance*

A	B
\$50	\$200

Serial schedule T_1, T_2 :

T₁ A += 100 B -= 100

$$T_2 \quad A^* = 1.06 \quad B^* = 1.06$$

A	B
\$159	\$106

Interleaved schedule A:

T₁ A += 100 B -= 100

T₂

A [*] = 1.06	B [*] = 1.06
-----------------------	-----------------------

A	B
\$159	\$106

Same
result!

Scheduling examples

Serial schedule T_1, T_2 :

T₂

A* = 1.06	B* = 1.06
-----------	-----------

Interleaved schedule B:

Diagram illustrating two parallel processes, T_1 and T_2 , each with a critical section. T_1 contains the operation $A += 100$, and T_2 contains the operation $B -= 100$. Both operations are enclosed in red boxes, indicating they are critical sections.

T₂

A* = 1.06	B* = 1.06
-----------	-----------

*Starting
Balance*

A	B
\$50	\$200

A	B
\$159	\$106

Different
result than
serial
 T_1, T_2 !

A	B
\$159	\$112

Scheduling examples

Serial schedule T_2, T_1 :

T_1

A += 100

B -= 100

T_2

A *= 1.06

B *= 1.06

Starting
Balance

A	B
\$50	\$200

A	B
\$153	\$112

Interleaved schedule B:

T_1

A += 100

B -= 100

T_2

A *= 1.06


B *= 1.06

A	B
\$159	\$112

Different
result than
serial
 T_2, T_1
ALSO!

Scheduling examples

Interleaved schedule B:




This schedule is different than ***any serial order!*** We say that it is **not serializable**

Scheduling Definitions


- A **serial schedule** is one that does not interleave the actions of different transactions
- A and B are **equivalent schedules** if, *for any database state*, the effect on DB of executing A **is identical to** the effect of executing B
- A **serializable schedule** is a schedule that is equivalent to ***some*** serial execution of the transactions.

The word “**some**” makes this definition powerful & tricky!

Serial Schedules




S1




S2


Interleaved Schedules




S3



S4



S5




S6

Serial Schedules	S1, S2
Serializable Schedules	S3, S4
Equivalent Schedules	<S1, S3> <S2, S4>
Non-serializable (Bad) Schedules	S5, S6


General DBMS model: Concurrency as Interleaving TXNs

Serial Schedule



Each action in the TXNs
*reads a value from global
memory and then writes
one back to it*

Interleaved Schedule



For our purposes, having TXNs
occur concurrently means
**interleaving their component
actions (R/W)**


We call the particular order
of interleaving a **schedule**

Anomalous/Bad data,
if we aren't careful

Classic Anomalies with Interleaved Execution

“Unrepeatable read”:

Example:




1. T_1 reads some data from A
2. T_2 writes to A
3. Then, T_1 reads from A again *and now gets a different / inconsistent value*

Occurring with / because of a ***RW conflict***

Classic Anomalies with Interleaved Execution

“Dirty read” / Reading uncommitted data:

Example:




1. T_1 writes some data to A
2. T_2 reads from A, then writes back to A & commits
3. T_1 then aborts- *now T_2 's result is based on an obsolete / inconsistent value*

Occurring with / because of a **WR conflict**

Classic Anomalies with Interleaved Execution

“Inconsistent read” / Reading partial commits:

Example:



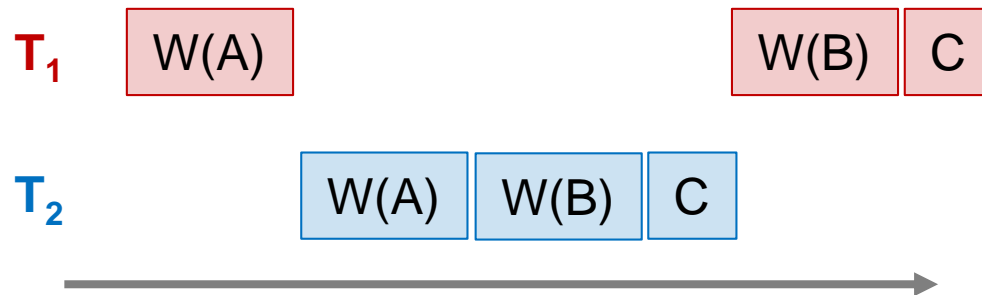
1. T_1 writes some data to A
2. T_2 reads from A and B , and then writes some value which depends on A & B
3. T_1 then writes to B - now T_2 's result is based on an incomplete commit

Again, occurring because of a **WR conflict**

Classic Anomalies with Interleaved Execution

Partially-lost update:

Example:



1. T_1 blind writes some data to A
2. T_2 blind writes to A and B
3. T_1 then blind writes to B; now we have T_2 's value for B and T_1 's value for A- **not equivalent to any serial schedule!**

Occurring because of a **WW conflict**

How to treat Conflicts carefully?

Conflict Types

Two actions **conflict** if they are part of different TXNs, involve the same variable, and at least one of them is a write

Thus, there are three types of conflicts:

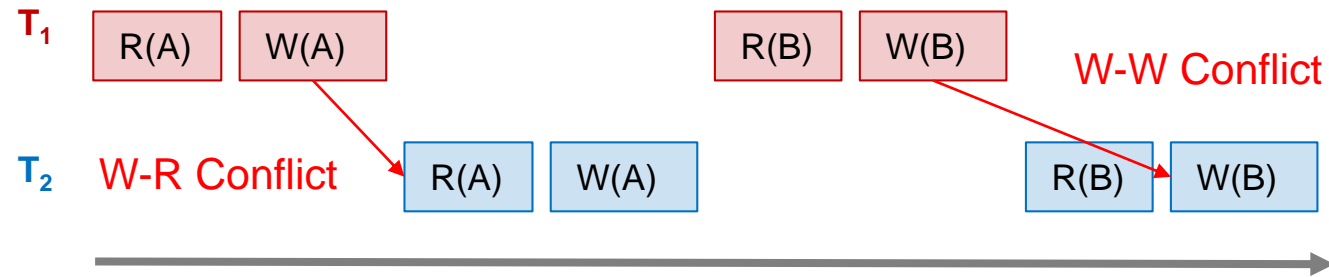
- Read-Write conflicts (RW)
- Write-Read conflicts (WR)
- Write-Write conflicts (WW)

Why no “RR Conflict”?

Note: **conflicts** happen often in many real world transactions. (E.g., two people trying to book an airline ticket)

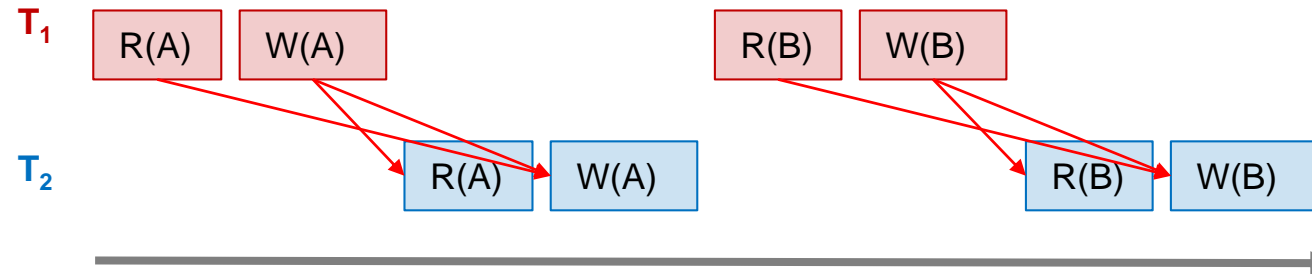
Conflicts

Two actions **conflict** if they are part of different TXNs, involve the same variable, and at least one of them is a write



Conflicts

Two actions **conflict** if they are part of different TXNs, involve the same variable, and at least one of them is a write



All "conflicts"!

Note: Conflicts vs. Anomalies

Conflicts are in both “good” and “bad” schedules
(they are a property of transactions)

Goal: Avoid Anomalies while interleaving transactions with conflicts!

- Do not create “bad” schedules where isolation and/or consistency is broken (i.e., Anomalies)

Conflict Serializability, Locking & Deadlock

Conflict Serializability

Two schedules are **conflict equivalent** if:

- They involve *the same actions of the same TXNs*
- Every *pair of conflicting actions* of two TXNs are *ordered in the same way*


Schedule S is **conflict serializable** if S is *conflict equivalent* to some serial schedule

Conflict serializable \Rightarrow serializable

So if we have conflict serializable, we have consistency & isolation!


Example “Good” vs. “bad” schedules

Serial Schedule:



Note that in the “bad” schedule, the **order of conflicting actions is different than the above (or any) serial schedule!**


Interleaved Schedules:



Conflict serializability provides us with an operative notion of “good” vs. “bad” schedules! “Bad” schedules create data Anomalies


The Conflict Graph

- Let's now consider looking at conflicts **at the TXN level**
- Consider a graph where the **nodes are TXNs**, and there is an edge from $T_i \rightarrow T_j$ **if any actions in T_i precede and conflict with any actions in T_j**




What can we say about “good” vs. “bad” conflict graphs?

Serial Schedule:




A bit complicated...

Interleaved Schedules:




What can we say about “good” vs. “bad” conflict graphs?

Serial Schedule:



Simple!

Interleaved Schedules:



Theorem: Schedule is **conflict serializable** if and only if its conflict graph is **acyclic**

Concurrency

Connection to conflict serializability

- In the conflict graph, a topological ordering of nodes corresponds to a **serial ordering of TXNs**
- Thus an **acyclic** conflict graph \rightarrow conflict serializable!

Theorem: Schedule is **conflict serializable** if and only if its conflict graph is **acyclic**

Two-Phase Locking (2PL)

- Consider *two-phase locking* - as a way to deal with concurrency
 - Guarantees conflict serializability
 - (if it completes- see upcoming...)
- Also (*conceptually*) straightforward to implement, and transparent to the user!

Two-phase Locking (2PL) Protocol

TXNs obtain:

- An **X (*exclusive*) lock** on object before **writing**.
 - If a TXN holds, no other TXN can get a lock (S or X) on that object.
- An **S (*shared*) lock** on object before **reading**
 - If a TXN holds, no other TXN can get an X lock on that object
- All locks held by a TXN are released when TXN completes.

Note: Terminology here- “exclusive”, “shared”- meant to be intuitive- no tricks!


Two Phase Locking Protocol

■ A transaction is said to follow Two Phase Locking protocol if Locking and Unlocking can be done in two phases

- Growing Phase: New locks on data items may be acquired but none can be released
- Shrinking Phase: Existing locks may be released but no new locks can be acquired

	T ₁	T ₂
1	LOCK-S(A)	
2		LOCK-S(A)
3	LOCK-X(B)	
4
5	UNLOCK(A)	
6		LOCK-X(C)
7	UNLOCK(B)	
8		UNLOCK(A)
9		UNLOCK(C)
10

Picture of 2-Phase Locking (2PL)



Strict 2PL

Theorem: Strict 2PL allows only schedules whose dependency graph is acyclic

Proof Intuition: If there is an edge $T_i \rightarrow T_j$ (i.e. T_i and T_j conflict) then T_j needs to wait until T_i is finished – so *cannot* have an edge $T_j \rightarrow T_i$

Therefore, Strict 2PL only allows conflict serializable
 \Rightarrow serializable schedules

Strict 2PL

If a schedule follows strict 2PL, it is conflict serializable...

- ...and thus serializable
- ...and we get isolation & consistency!

Popular implementation

- Simple, produces subset of *all* conflict serializable schedules
- One key, subtle problem (next)

HW Assignment #2

■ Create a simple table

```
1 create table Worker(  
2 WorkerID varchar(20),  
3 WorkerName varchar(255),  
4 WorkerJob varchar(255),  
5 CONSTRAINT PK_Worker PRIMARY KEY (WorkerID)  
6 );
```

```
1 insert into worker values ('1','John', 'nurse');  
2 insert into worker values ('2','Grace', 'farmer');  
3 insert into worker values ('3','Smith', 'doctor');
```

■ Do locking test for a specific record (i.e., one tuple) using two different transactions

- To obtain a lock for a specific record, you can use a command “select .. from ... where.... for update”. Then, the tuples satisfying the conditions will be locked. To use two different transactions and check their results, open two SQL Plus windows.

■ Test the following cases

- [Case1] In the first transaction, lock a record; In the second transaction, try to access the same record
 - **Capture the screen where the second transaction is blocked while presenting the used queries**
- [Case2] To release the lock in the first transaction, execute “commit” where the previous lock has been obtained
 - **Capture the screen where the second transaction is unblocked**
- [Case3] Try to test obtaining the lock for a different record from the second transaction while the first transaction has a lock for a specific record as in Case1
 - **Capture the screen where the second transaction can access the data**

■ Submit three screen shots above

Example: Deadlock Detection



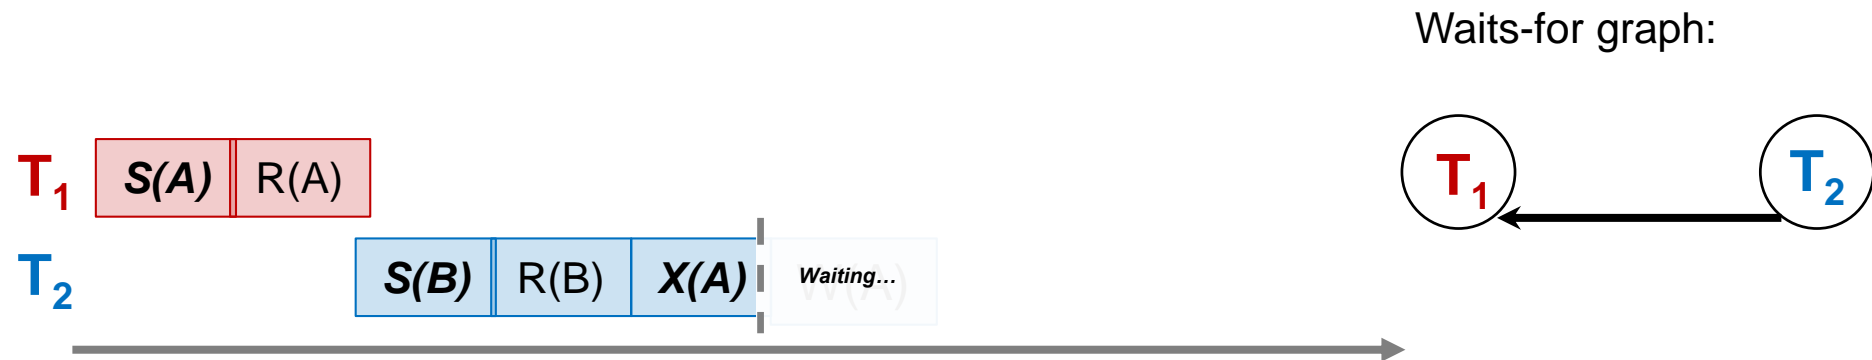
First, T_1 requests a shared lock on A to read from it

Deadlock Detection: Example




Next, T_2 requests a shared lock on B to read from it

Deadlock Detection: Example




T_2 then requests an exclusive lock on A to write to it- **now T_2 is waiting on T_1 ...**

Deadlock Detection: Example



Finally, T_1 requests an exclusive lock on B to write to it- **now T_1 is waiting on T_2 ... DEADLOCK!**

Waits-for graph:



Cycle =
DEADLOCK

Deadlocks

Deadlock: Cycle of transactions waiting for locks to be released by each other.

Two ways of dealing with deadlocks:

- Deadlock prevention
- Deadlock detection

Deadlock Detection

Create the **waits-for graph**:

- Nodes are transactions
- There is an edge from $T_i \rightarrow T_j$ if T_i is *waiting for T_j to release a lock*

Periodically check for (***and break***) cycles in the waits-for graph

Summary

- Concurrency achieved by **interleaving TXNs** such that **isolation & consistency** are maintained
 - We formalized a notion of **serializability** that captured such a “good” interleaving schedule
- We defined **conflict serializability**
- **Locking** allows only conflict serializable schedules
 - If the schedule completes... (it may deadlock!)

Transactions Summary

Why study Transactions?

Good programming model for parallel applications on shared data !

Atomic
Consistent
Isolation
Durable

Design choices?

- Write update Logs (e.g., WAL logs)
- Serial? Parallel, interleaved and serializable?

